

CAPTURE-RECAPTURE WITH COVARIATES: THE BAYESIAN LOGISTIC CAPTURE-RECAPTURE MODEL WITH EXTENSIONS

Robert Edward Granger

Submitted to the faculty of the Univesity Graduate School
in partial fulfillment of the requirements
for the degree
Doctor of Philosophy
in the Department of Statistics,
Indiana University
May 2024 (set this to month in which all requirements are fulfilled in title.sty)

Accepted by the Graduate Faculty, Indiana University, in partial fulfillment of the requirements for the degree of Doctor of Philosophy.

Doctoral Committee:

Daniel Manrique-Vallier, Ph.D.

Julia Fukuyama, Ph.D.

Amanda Mejia, Ph.D.

Roni Khardon, Ph.D.

May 1, 2024 (change this to defense date, it's in title.sty)

Robert Edward Granger

CAPTURE-RECAPTURE WITH COVARIATES: THE BAYESIAN LOGISTIC
CAPTURE-RECAPTURE MODEL WITH EXTENSIONS

Capture-recapture refers to a series of methods that are used to estimate the size of a population from at least two incomplete, matched lists. With human populations, heterogeneity in the capture probability is a serious concern as it can impact the estimation process. This proposal outlines a framework for incorporating heterogeneity through known covariates into the estimation process and develops an extensible model under this framework which we call the Bayesian Logistic Regression model for Capture-Recapture (BLRCR). The posterior distribution is estimated using the complete likelihood through a Markov Chain Monte Carlo (MCMC) algorithm. Through a careful examination of the problem and a series of demonstrations, three important areas of future research become apparent for the dissertation project. First, we should consider the impact of the specification of the covariate distribution and explore additional non-parametric solutions. Second, we should consider the unobserved heterogeneity and specifically how to implement a procedure that accounts for this violation in the conditional independence assumption. Third, we should consider the impact of variable selection.

Contents

1	Introduction	1
2	Literature Review	5
2.1	Early Approaches to List Dependency and Homogeneity	6
2.2	Capture-Recapture with Covariates	8
2.3	Bayesian Methods for Capture-Recapture	10
3	The Bayesian Capture-Recapture Model with Covariates	12
3.1	Background and Preliminaries	14
3.1.1	The Classical Capture-Recapture Approach	14
3.1.2	Capture-Recapture with Covariates	15
3.2	The Bayesian Logistic Regression Capture-Recapture Model	17
3.2.1	The BLRCR Model	17
3.2.2	Estimation of the BLRCR	19
3.3	Simulation Analysis	21
3.3.1	Simulations with Varying Levels of List Dependency	24
3.3.2	Simulations with Different Sized Populations	26
4	Selecting a Distribution for the Covariates	29
4.1	Specifying a Distribution for the Covariates	30
4.2	Conditional Likelihood	31
4.3	Nonparametric Distributions	32
4.3.1	Dirichlet Process Mixture of Normal Distributions	32

4.3.2	The Bayesian Bootstrap	35
4.4	Simulations using Different Covariate Distributions	37
5	Conditional Independence and Unobserved Heterogeneity	43
5.1	Modelling the Unobserved Heterogeneity with Latent Classes	44
5.2	Updating the Estimation	45
5.3	Simulations with Unobserved Heterogeneity	46
6	Model Selection	51
6.1	Model Comparison via the Bayes Factor	51
6.2	Regularization and Prior Specifications	55
7	Example 1 (probably superhero)	57
8	Example 2 (find something real even if we don't know the answer)	58
9	Conclusion	59
9.1	Unobserved Heterogeneity	59
9.2	Modelling the Covariate Distribution	59
9.3	Covariate Selection	60
10	Appendix A: Conditional Maximum Likelihood Estimation	61
11	Appendix B: LCMCR Model	63
12	References	65
	Curriculum Vitae	

1 Introduction

Capture-Recapture (CR) refers to a series of methods that are used to estimate the size of a population from at least two incomplete, matched lists. While one could simply count the number of unique names between the lists, this would provide only a lower bound estimate of the population size. Presumably, there are individuals who were not captured in any of the lists. CR methods use assumptions along with patterns in the data to create estimates for how many individuals were missed and therefore inherently the size of the total population.

Most early applications of capture-recapture, also called mark and recapture, were in the context of animal populations where a number of animals were initially captured, marked or tagged, and then released. This was followed up with another capturing occasion where animals were captured and checked for a marking. A CR dataset is then a listing of each unique animal's capture pattern after all sampling occasions have taken place. CR methods have been used in a wide array of additional subject areas including ecology ([Henderson and Southwood, 2016](#)), epidemiology ([Goldberg and Wittes, 1978](#); [Baker, 1990](#)), U.S. Census adjustments ([Darroch et al., 1993](#)), and estimating the number of pages on the internet ([Lawrence and Giles, 1998](#); [Khabisa and Giles, 2014](#)). Of particular interest to this dissertation is the application of this methodology to human populations, which are sometimes referred to as multiple systems estimation (MSE). Lists of subjects may be collected by various institutions including hospitals, government agencies, or other non-governmental organizations but are left incomplete due to lack of resources, mistakes/omissions, or just the inability to obtain such information. These lists can then be combined and constructed to form a dataset for CR. For example, [Zwane and van der Heijden \(2005\)](#) attempts to determine the number of Dutch children born with a neural tube defect (NTD) in the year 2000 by analyzing three national databases tracking these occurrences. For multiple reasons including risk assessment of the pregnancy and

omissions, not all children with NTDs are reported to each database. Another example is the counting of casualties from various conflicts including Colombia (Guberek et al., 2010; Manrique-Vallier, 2016), Guatemala (Ball et al., 2000; Ball and Price, 2018), Kosovo (Ball et al., 2002; Manrique-Vallier, 2016), Peru (Ball et al., 2003), and many others. Casualty counts by a reporting agency may only be partially collected because of non-cooperation from victims due to lack of trust, danger posed to data collectors, and/or destroyed infrastructure (Manrique-Vallier et al., 2020).

CR methods often come with many assumptions including a closed population, independence between lists, and homogeneity in the capture probability between individuals. In particular, the aim of the dissertation is to extend methods that target the homogeneity assumption, the requirement that all individuals have the same probability of capture on each list regardless of any personal attributes. While this may be a palatable assumption for some animal populations, this is unlikely with human populations. Different characteristics like age or social status may influence the probability of their capture. If data is present that can fully explain the heterogeneity, one could incorporate this by stratifying before applying their CR method of choice (Sekar and Deming, 1949; Manrique-Vallier et al., 2019). This is only possible if the covariate is discrete or can be made discrete. It can also lead to issues of sparsity. An alternative is to include the covariates in a model like regression (Alho, 1990; Baker, 1990; Bonner and Schofield, 2014; King et al., 2016). Regardless, this only addresses the observable heterogeneity and ignores heterogeneity that may exist but covariate information is missing or is not fully adequate. Some approaches have been proposed to account for unobservable heterogeneity by adapting the Rasch model (Darroch et al., 1993) or latent class memberships (Manrique-Vallier, 2016).

This dissertation develops a fully Bayesian procedure for the CR problem that has a number of desirable characteristics:

- 1) allows for multiple lists;
- 2) is resistant to sparsity;
- 3) can be informed through the use of priors;
- 4) uses discrete and continuous covariates to account for observable heterogeneity;
- 5) can also account for unobservable heterogeneity;
- 6) extensible;

Estimating the size of the population tends to be highly susceptible to the structure of the modeling when dealing with CR methods. Therefore, implementing covariates in an improper way or not adequately accounting for the unobservable heterogeneity can bias the inference. As a result, we propose developing a method that is fully Bayesian approach and uses the full likelihood as suggested by [Pollock \(2002\)](#). We begin by developing a framework for CR that incorporates covariates, a framework that for all intents and purposes is identical to the complete likelihood specification in ([King et al., 2016](#)). Using this framework, we specify a model that uses conditionally independent logistic regressions which we term the Bayesian Logistic Regression Capture-Recapture (BLRCR) model.

While the BLRCR model incorporates individual covariates and thus accounts for observable heterogeneity, it has a few shortcomings:

- 1) Since covariates are missing at least for the unobserved individuals, the use of a full likelihood approach requires the specification of a distribution for the covariates. As we will see in [Section 4](#), misspecifying the distribution can lead to incorrect inference. Like previous approaches, we apply a normal distribution; however, we propose a new method of using non-parametric approaches to the covariates such as an infinite mixture of normal distributions.

- 2) Unobservable heterogeneity can bias the inference, and hence we need a way to account for it. From our perspective, we view the idea of unobservable heterogeneity as an issue of omitted covariates. We borrow the idea of individuals belonging to latent classes impacts the probability of capture([Manrique-Vallier, 2016](#)), but implement it in the regression procedure through an additional intercept term.
- 3) Determining which covariates to include and how each covariate should be included in the regression will ultimately impact the inference. There are multiple ways to account for this issue, but we explore it through the lens of variable selection. We investigate the difficulty of computing the Bayes Factor and ultimately suggest a regularization technique.

The dissertation proceeds as follows: Section [2](#) provides a brief history of the capture-recapture problem while reviewing literature relevant to Bayesian and covariate extensions. Section [3](#) begins with preliminary details, provides a model using conditional independent Bayesian Logistic Regressions, and approximates the posterior distribution using a Markov Chain Monte Carlo (MCMC) estimation procedure. Section [4](#) explores different covariate distributions along with nonparametric specifications. Section [5](#) introduces hidden covariates that allow for a breaking of the conditional independence assumption. Section [6](#) investigates the problem of model selection. Section [7](#) looks at Example 1. Section [8](#) looks at Example 2. Section [9](#) concludes with a summary of the dissertation and future directions of work in this area.

2 Literature Review

Capture-Recapture (CR) refers to a series of methods that are used to estimate the size of a population from at least two incomplete, matched lists. This naming comes from the process of first capturing/marking a number of organisms (capture) and following up with a second capturing occasion (recapture). One then estimates the population size, N , of the organism by comparing the number of organisms marked in each sample, n_1 and n_2 , and the organisms that were matched in both samples, n_{12} . Intuitively, the proportion of the total population that is initially captured/marked, n_1/N , should be equal to the proportion of captured/marked individuals that are recaptured, n_{12}/n_2 . Of course, this only makes sense if we believe the lists are independent. Rearranging and solving for N leads to the formula:

$$\hat{N} = \frac{n_1 \cdot n_2}{n_{12}}.$$

Although this estimator was used as early as 1783 by the mathematician Pierre-Simon Laplace for the purpose of estimating the French population ([Schaefer, 1951](#)), this estimator is commonly referred to as the Lincoln-Petersen estimator after its early use by Frederick C. Lincoln on migratory waterfowl ([Lincoln, 1930](#)) and C.G. Johannes Petersen on the fish species plaice ([Petersen, 1895](#)). The Lincoln-Petersen estimator is relatively simple to compute, but is limited to only two lists with the assumption of a closed population. A closed population means the population does not change throughout the sampling occasions. This estimator also makes two other intricately related assumptions:

- 1) independence between lists,
- 2) homogeneity in capture probabilities between individuals.

The first assumption of independence between lists means the probability an individual is captured on one list does not depend on the probability that individual is captured on

another list. The second assumption means the probability of any two individuals being captured on a list is the same. When these assumptions are violated, it can lead to bias in our population estimates. Addressing violations of these assumptions has a long history and continues to be looked into today.

2.1 Early Approaches to List Dependency and Homogeneity

Capture-recapture solutions with multiple lists appear at least as early as the 1920s ([Geiger and Werner, 1924](#); [Schnabel, 1938](#); [Darroch, 1958](#)). All use a model for estimating the size of a population from multiple lists with list independence. If this assumption were to hold, then the probability an individual appears on one list would not impact the probability that the individual appears on another list. Perhaps for carefully controlled experiments, list independence can be assumed, but for most real applications, this is unlikely. For example with human rights data, this may be violated as psychological research indicates survivors of human rights abuses can benefit in their mental health by sharing their testimony. These individuals may in turn be encouraged to share their experiences to multiple outlets. Furthermore, there is evidence with the Guatemalan Civil War that several popular movements groups encouraged their social bases to provide testimony of human rights abuses to all of three of the data collecting organizations ([Ball et al., 2000](#)). These conditions lead to positive list dependence and therefore an underestimate of the population size. On the other hand, there is also the potential that certain groups of individuals may feel uneasy sharing their testimony with non-aligned data collection groups, leading to negative dependence and an overestimate of the population size. For example, religious members of the Catholic faith were more apt to give their testimony to the Catholic researchers than they were to the political left, non-governmental organizations ([Manrique-Vallier et al., 2020](#)).

In the early 1970s, two approaches were developed to address list dependency. [Fienberg](#)

(1972) addressed the issue of list dependency directly by modeling interactions between lists using hierarchical log linear models under a multinomial sampling scheme. On the other hand, [Sanathanan \(1972\)](#) also used a multinomial sampling scheme but instead of modeling list dependence directly, she modeled it as unobserved heterogeneity in the individuals under a conditional likelihood approach. As pointed out in several places ([Darroch et al., 1993](#); [Fienberg et al., 1999](#); [Manrique-Vallier et al., 2020](#)), list dependency and heterogeneity between individuals are intricately related such that list dependency can be viewed as a result of heterogeneity between individuals. This leads to the second assumption, homogeneity, which refers to all individuals having the same capture probability within a list. For a variety of reasons, we may expect certain individuals to be more likely to be captured because of characteristics like geographic area, age, or their status within the community.

This heterogeneity among individuals may be unobserved, observed through covariates, or a combination. In the 1990s, various methods were proposed to account for unobserved heterogeneity using adaptations to log-linear and mixture models ([Darroch et al., 1993](#); [Agresti, 1994](#); [Pledger, 2000](#)). Later, Bayesian methods based on latent class models were presented that allowed for a more flexible approach to heterogeneity ([Fienberg et al., 1999](#); [Manrique-Vallier and Fienberg, 2008](#); [Manrique-Vallier, 2016](#)). These methods seek out patterns in the capture histories in order to infer for unobserved heterogeneity. When the researcher has accompanying discrete covariate data, a suggestion noted by [Darroch et al. \(1993\)](#), is to first stratify the data based on the observable heterogeneity ([Sekar and Deming, 1949](#)), and then perform the desired method within that stratification. An example of this this implementation can be found in [Manrique-Vallier et al. \(2019\)](#), where the authors estimate the number of casualties in the internal conflict in Peru (1980-2000) based on geographic stratifications. Stratification becomes problematic in the presence of continuous covariates. When continuous covariates are present, the researcher must make difficult decisions on how to discretize the covariates which leads to assigning potentially

arbitrary cutoffs. If your stratification scheme results in too many strata, the number of observed individuals within each strata may be small, resulting in loss of inferential power and identifiability. These issues can arise even without continuous covariates if there are many discrete variables with many unique levels.

2.2 Capture-Recapture with Covariates

An alternative to stratification is to use the covariates within the modeling process like with regression. [Pollock et al. \(1984\)](#) suggested the use of logistic regression with environmental and individual level covariates. [Huggins \(1989\)](#); [Alho \(1990\)](#) independently proposed a similar but slightly different approach that uses conditional maximum likelihood logistic regression through a two step process. First, the coefficients for the logistic regression model are obtained by regressing whether the individual was captured against the covariates under the condition the individual is observed at least once. If we perform the regression without taking into account that only observed individuals appear in the dataset, i.e., some individuals are missing, the coefficients become biased. After obtaining the coefficients, we fit the probability that each individual is captured on each list, and the probability an individual is missing, θ_i , is computed. The second step uses the estimated probability of missing by plugging it into the [Horvitz and Thompson \(1952\)](#) estimator,

$$\hat{N} = \sum_i^n \frac{1}{\theta_i},$$

to obtain an estimate for N . Intuitively, the estimator inverts the probability of missing to account for the number of similar individuals that would have been unobserved. This intuition implies that the estimator will not perform well when the covariates for the observed individuals vary significantly from the unobserved individuals. Nevertheless, [Alho \(1990\)](#) shows this estimator to be consistent with large samples and develops an asymptotic approach to estimating the variance.

In these models, the probability an individual is captured on each list is conditionally independent of the other lists conditional on the covariates. While covariates may explain some of the heterogeneity, it is unlikely that all heterogeneity between individuals can be explained in such a fashion and that some list dependency would not remain. [Zwane and van der Heijden \(2005\)](#) use multinomial logistic regression with a design matrix intended to directly model the dependence between the lists. One could theoretically incorporate additional rows into the design matrix to account for unobserved heterogeneity as in [Darroch et al. \(1993\)](#).

Since the covariates will be missing for any of the unobserved individuals, any approach that uses the complete likelihood would require the covariate distribution to be specified. All of the previous mentioned methods sidestep this need by maximizing the conditional likelihood instead of the full likelihood. [Stoklosa et al. \(2011\)](#) takes this a step further by comparing a partial likelihood approach, i.e, the number of recaptures after the first capture, with the conditional approach. They find a loss of efficiency but argue that it allows more flexibility in modeling. Later, [Yee et al. \(2015\)](#) presented a simple way to implement conditional likelihood methods using the **VGAM** package in the programming software **R**. There is some early usage of a full likelihood approach appearing in work that combines line transect theory with capture-recapture for population estimation ([Alpizar-Jara and Pollock, 1996](#); [Borchers et al., 1998](#)). Since the full likelihood approach requires specifying a distribution for the covariates, this presents the difficulty of potentially estimating additional parameters regarding the covariate distribution and/or integrating out the covariates to obtain an estimate for N . While more challenging, [Pollock \(2002\)](#) speculates the full likelihood approach could result in better precision and suggests a possible solution may lie with Bayesian methods.

2.3 Bayesian Methods for Capture-Recapture

Perhaps the earliest paper to utilize Bayesian methods is [Roberts \(1967\)](#) in the evaluation of “stopping rules” with exactly two samples. The “stopping rule” refers to how the sample of individuals from the population will be collected, and in particular, what criterion must be met before concluding the sampling occasion. For example, a researcher may choose a fixed sample size and will stop sampling subjects when this sample size is reached. On the other hand, the researcher may attempt to sample every subject within a fixed time period. Depending on the chosen stopping rule, the likelihoods can differ and thus the estimator. See [Chapman \(1954\)](#); [Darroch \(1958\)](#) for a non-Bayesian look at the problem. These different stopping rules were once again considered in the Bayesian approaches of [Castledine \(1981\)](#) and extended to include more than two lists. Further, the authors assume the probability of capture is constant between individuals but allowed to vary across time (see model M_t in [Otis et al. \(1978\)](#)). The impact of various prior specifications is explored including the use of the Beta prior for the probability of capture and the Jeffreys prior ([Jeffreys, 1967](#)) for N . Under similar modelling, estimates of the posterior distribution were obtained using empirical Bayes and Bayes empirical Bayes ([Smith, 1991](#)), and Markov Chain Monte Carlo (MCMC), Gibbs sampling, ([George, 1992](#)) methodology.

Bayesian approaches to unobserved heterogeneity were first introduced to the multiple-recapture problem in [Fienberg et al. \(1999\)](#). The authors use a fully Bayesian hierarchical approach through the use the Rasch model and develop a MCMC procedure for obtaining samples from the posterior. [Manrique-Vallier and Fienberg \(2008\)](#) also employ a fully Bayesian hierarchical approach but use a Grade of Membership model ([Woodbury et al., 1978](#)) to account for the unobserved heterogeneity. Individuals are soft or partially clustered into classes, i.e. mixed membership, with each class having an estimated probability of capture per list. A simplification to the partial membership approach is to use the Latent Class Model (LCM) where individuals are hard clustered into

groups. (Manrique-Vallier, 2016) uses a Bayesian Nonparametric Latent Class Model (NPLCM) by using a potentially infinite number of hidden classes through a stick-breaking prior (Dunson and Park, 2008). Posterior samples are obtained via a Gibbs sampler.

Naturally, when deriving a posterior distribution, one would use the full or “complete” likelihood. This has the benefit of 1) being the “correct” way to compute the posterior based on a specified generative process and 2) allows us to extend our model without fear of compounding any errors from modifications/approximations. Unfortunately, the use of a full likelihood approach in the presence of covariates poses a particularly challenging problem as a distribution for the covariates must be specified and the posterior often becomes intractable. The early approach of using the conditional likelihood instead of the full likelihood avoids this problem but recently several other solutions have been presented. These include the use of a reversible jump MCMC (King and Brooks, 2008), data augmentation (Royle et al., 2007; Royle, 2009), and replacing the full likelihood with a “semi-complete” data likelihood approach (King et al., 2016). The semi-complete likelihood uses the complete likelihood for the observed individuals and a marginal likelihood for the unobserved individuals. Both approaches still require estimating the probability of missing, which is a difficult integral. While (King et al., 2016) evaluate the integral using Gauss-Hermite quadrature, Bonner and Schofield (2014) estimates the probability of missing with a Monte Carlo within MCMC step. The efficiency of this approach depends on the size of the Monte Carlo simulated covariate distribution with a larger sample giving a better approximation, but at the cost of reducing the efficiency of the algorithm. It should be noted that whether we use a full likelihood or even a semi-complete likelihood, a covariate distribution must be specified. The examples provided often assume a normal distribution, but Royle (2009) does consider alternative distributions. These alternative distributions do reveal some sensitivity to the choice of covariate distribution, which will be explored further in Section 4.

3 The Bayesian Capture-Recapture Model with Covariates

In this section, we develop a Bayesian procedure for the multiple-list capture-recapture (CR) problem with covariates.

Perhaps the most ubiquitous technique in solving the CR problem with multiple-lists is through the frequentist technique of log-linear models (Fienberg, 1972). Using multiple lists ($J \geq 2$) can be beneficial in that we have more detailed patterns, which can lead to better inference. The downside is that as the number of lists increases, the number of possible patterns grows exponentially. For example, with just two lists, there are only four possible patterns: 1) the individual shows up on both lists, 2) the individual shows up only on list 1, 3) the individual shows up only on list 2, or 4) the individual shows up on neither list. The number of possible patterns can be calculated as 2^J , so a dataset with 30 lists would have over a billion possible patterns. Not only could this become computationally expensive, but it also leads to issues of sparsity, i.e., many of the potential patterns will not appear in the data. In Manrique-Vallier (2016), they analyze a dataset about killings in Casanare, Colombia which contains 15 lists. They note that only 70 of the potential $2^{15} = 32,768$ capture patterns are present and as a result were unable to successfully compute a solution using log-linear models.

Even data with much fewer lists can still have issues of sparsity. To get around this issue, we present the Bayesian Logistic Regression Capture-Recapture model that relies on an assumption of conditional independence. First, since the method is Bayesian, we have the added benefit of allowing the practitioner to insert prior belief or knowledge into the estimation procedure. Bayesian methods also are known to assist with sparsity in that they essentially “create” data to fill in the sparse areas. In addition, the assumption of conditional independence, while strong, allows for a reduction in the complexity and sparsity issue. The conditional independence assumption states that given the covariates, the probability of capture on one list is unaffected by another list.

Unlike most popular existing methods, the method presented here introduces the ability to insert covariates to guide in the estimation of the capture probability. Instead of relying on stratification that essentially only allows for discrete covariates, we implement conditionally independent logistic regressions on each of the lists to determine the probability of capture (or inversely, the probability of non-capture). Typically with regression, the distribution of the covariate is irrelevant so it does not matter whether it is discrete or continuous.

Unfortunately, since we’re dealing with what amounts to a missing data problem, but the missing data is not missing at random ([Rubin, 1976](#)), the distribution of the covariate matters (see [Section 4](#)). Nevertheless, while we must be careful how we model the distribution of the covariate, any type of covariate, continuous or discrete, can be used.

Finally, a key concern of any capture-recapture framework is that the estimation is highly dependent on how the model is structured. Different situations may present important characteristics of the problem that need to be carefully considered by the practitioner. With that in mind, the desire is to create a methodology that is extensible. Through the use of the full likelihood and a data augmentation approach, we are able to add extensions more readily including more robust ways of handling heterogeneity. In particular, we may be concerned with both “observable” heterogeneity, i.e. the heterogeneity that can be detected through the use of covariates, but we may also be concerned with “unobservable” heterogeneity, i.e. the heterogeneity that persists but for which no covariate information exists. One such method, presented in [Section 5](#), allows for the modeling of this unobservable heterogeneity through the use of latent classes. This assumption also happens to break the perhaps unpalatable assumption of conditional independence mentioned before. Similarly, the covariates that are present may be numerous or may affect the estimation procedure in non-linear ways, so we extend the model to include a variable selection procedure (see [Section 6](#)). Additional extensions may be warranted given the problem, but the estimation procedure that uses data augmentation allows for extensions to be added relatively easily.

To summarize, we begin with a Bayesian framework for the capture-recapture problem with covariates (see subsection 3.1). Next, we propose a specific model following this procedure which we call the Bayesian Logistic Regression Capture-Recapture (BLRCR) model and estimate the posterior through a Markov Chain Monte Carlo (MCMC) algorithm (see subsection 3.2). We then explore the properties of this model using a series of simulations in Section 3.3.

3.1 Background and Preliminaries

3.1.1 The Classical Capture-Recapture Approach

We begin with the multinomial multiple-recapture framework first presented in Darroch (1958) and utilized by numerous past and more recent works (Sandland and Cormack, 1984). This section summarizes this framework, heavily relying on the notation of Manrique-Vallier (2016). The objective of this capture-recapture framework is to estimate the unknown size, N , of a population of individuals assuming the population size remains unchanged throughout the capturing occasions, i.e. a closed population, and that the captured individuals can be matched perfectly. Individuals are captured (sampled) through the use of $J \geq 2$ lists. If individual, i , is captured on list j , then $y_{ij} = 1$ with $y_{ij} = 0$ otherwise. When aggregated across lists, we refer to these values as capture vectors, $\mathbf{y}_i = (y_{i1}, \dots, y_{iJ})$. If an individual is not identified on any of the J lists, $\mathbf{y}_i = \mathbf{0} \equiv (0, \dots, 0)$, then that individual is considered “unobserved” or “missing”. The number of unobserved individuals, n_0 , plus the number of individuals that are captured on at least one list, n , is equal to the size of the population, i.e, $N = n_0 + n$.

Each individual’s capture vector, \mathbf{y}_i , is generated from probability distribution, $f(\mathbf{y}|\theta)$ for all $i = 1, \dots, N$. Reordering the individuals such that the observed individuals are $1, \dots, n$ and the unobserved individuals are $n + 1, \dots, N$ leads to the following joint likelihood

$$p(\mathcal{Y}_{obs}|N, \boldsymbol{\theta}) = \binom{N}{n} f(\mathbf{0}|\boldsymbol{\theta})^{N-n} \prod_{i=1}^n f(\mathbf{y}_i|\boldsymbol{\theta}) I(N \geq n), \quad (1)$$

where $\mathcal{Y}_{obs} = (\mathbf{y}_1, \dots, \mathbf{y}_n)$, the observed capture vectors. For the classical CR problem, the only data is the observed capture histories, \mathcal{Y}_{obs} . The parameters of interest are N and $\boldsymbol{\theta}$. We place a prior distribution, $p(N, \boldsymbol{\theta})$, on the parameters with the objective of computing the posterior distribution,

$$p(N, \boldsymbol{\theta}|\mathcal{Y}_{obs}) \propto p(\mathcal{Y}_{obs}|N, \boldsymbol{\theta})p(N, \boldsymbol{\theta}). \quad (2)$$

3.1.2 Capture-Recapture with Covariates

We expand the framework from the previous section by allowing each individual's capture probability to be dependent on a matrix of covariates, \mathcal{X} . Let x_{ih} represent the value of covariate $h \in 1, \dots, H$ for individual $i \in 1, \dots, N$. Since some individuals are not captured, the covariate information for these individuals is lost. Let $\mathcal{X}_{obs} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$ be the covariate data on the $1, \dots, n$ individuals that are observed and let $\mathcal{X}_{mis} = (\mathbf{x}_{n+1}, \dots, \mathbf{x}_N)$ be the covariate data on the $n+1, \dots, N$ individuals that are not observed.

We update the joint likelihood in [Equation 1](#) to include these covariates,

$$\begin{aligned} p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|N, \theta, \boldsymbol{\phi}, \mathcal{X}_{mis}) &= p(\mathcal{Y}_{obs}|N, \theta, \boldsymbol{\phi}, \mathcal{X}_{obs}, \mathcal{X}_{mis})p(\mathcal{X}_{obs}|\boldsymbol{\phi}) \\ &= \binom{N}{n} \prod_{i=1}^n f(\mathbf{y}_i|\boldsymbol{\theta}, \mathcal{X}_{obs}) \prod_{i=n+1}^N f(\mathbf{0}|\boldsymbol{\theta}, \mathcal{X}_{mis}) I(N \geq n) \cdot g(\mathcal{X}_{obs}|\boldsymbol{\phi}). \end{aligned} \quad (3)$$

Once again, in order to complete the Bayesian model, priors must be assigned to the unknown parameters $\boldsymbol{\theta}$ and N ; however, we also must consider the distribution of the observed and missing covariates along with the parameter(s) governing their distribution,

ϕ . The joint posterior can be written as,

$$\begin{aligned} p(N, \boldsymbol{\theta}, \phi, \mathcal{X}_{mis} | \mathcal{Y}_{obs}, \mathcal{X}_{obs}) &\propto p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | N, \boldsymbol{\theta}, \phi, \mathcal{X}_{mis}) p(N, \boldsymbol{\theta}, \phi, \mathcal{X}_{mis}) \\ &= p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | N, \boldsymbol{\theta}, \phi, \mathcal{X}_{mis}) p(\mathcal{X}_{mis} | N, \phi) p(N, \boldsymbol{\theta}, \phi) \end{aligned} \quad (4)$$

Unfortunately, Equation 4 does not allow us to simply compute the posterior by conditioning on the covariates as in classical regression. Let $\mathcal{Y} = [\mathcal{Y}_{obs}, \mathcal{Y}_{mis}]_{N \times J}$ and $\mathcal{X} = [\mathcal{X}_{obs}, \mathcal{X}_{mis}]_{N \times H}$ represent the complete data; however, the unobserved portions of each of these matrices are fundamentally different. Conceptually, $\mathcal{Y} = [\mathcal{Y}_{obs}, \mathcal{Y}_{mis}]$ is decomposed into observed and unobserved components; however, $\mathbf{y}_i = \mathbf{0}$ for all $\mathbf{y}_i \in \mathcal{Y}_{mis}$, i.e., the values of each row are known. On the other hand, we do not know the values of the missing covariates, \mathcal{X}_{mis} . This complicates our ability to compute the posterior, which could have been a simple case of conditioning on the covariates as in classical regression. Because of the missing covariates, we specify a distribution for all covariates,

$$\mathbf{x}_i \stackrel{iid}{\sim} g(\phi). \quad (5)$$

To further complicate matters, the observed and unobserved covariates almost surely do not follow the same distribution as they are not missing at random. While $\mathbf{x}_i \in \mathcal{X}_{obs}$ or $\mathbf{x}_i \in \mathcal{X}_{mis}$ is defined through the observational status of its corresponding \mathbf{y}_i , all \mathbf{x}_i in each set is drawn independently from Equation 5. Nevertheless, it is not the case that $\mathbf{x}_i | \mathbf{y}_i \neq \mathbf{0}$ or $\mathbf{x}_i | \mathbf{y}_i = \mathbf{0}$ will necessarily have this same distribution. In other words, if we desire samples of our missing covariates, it would be incorrect to sample simply from $g(\phi)$, but instead these values must be sampled conditioned on the individual being missing.

3.2 The Bayesian Logistic Regression Capture-Recapture Model

3.2.1 The BLRCR Model

Using the framework from subsection 3.1.2, we implement the Bayesian Logistic Regression Capture-Recapture (BLRCR) model which uses independent logistic regressions to estimate the capture probabilities on each list. Suppose y_{ij} and x_{ih} are generated in the following way:

$$y_{ij}|\mathbf{x}_i \stackrel{ind}{\sim} \text{Bernoulli}(\lambda_{ij}(\mathbf{x}_i)) \quad \text{for } i = 1, \dots, N \text{ and } j = 1, \dots, J \quad (6)$$

$$\mathbf{x}_i \stackrel{iid}{\sim} \mathbf{g}(\boldsymbol{\phi}) \quad \text{for } i = 1, \dots, N, \quad (7)$$

where

$$\lambda_{ij}(\mathbf{x}_i) = \sigma(\mathbf{x}_i^T \boldsymbol{\beta}_j) = \frac{1}{1 + e^{-\mathbf{x}_i^T \boldsymbol{\beta}_j}}, \quad (8)$$

and $\mathbf{g}(\boldsymbol{\phi})$ is the distribution of the H covariates. The capture probability that individual i appears on list j is equal to λ_{ij} , and this value can be calculated with a nonstochastic transformation of \mathbf{x}_i and $\boldsymbol{\beta}_j$ through the sigmoid function (see Equation 8). In the linear term, $\mathbf{x}_i^T \boldsymbol{\beta}_j$, a different set of $\boldsymbol{\beta}_j$ covariates are used to determine each λ_{ij} implying a total of $k \times (h + 1)$ coefficients with the inclusion of an intercept. Notice this setup implies each individual's capture pattern is independent conditional on the covariates, i.e,

$$p(\mathbf{y}_i|\mathbf{x}_i) = \prod_{j=1}^J p(y_{ij}|\mathbf{x}_i, \boldsymbol{\theta}).$$

King et al. (2016) proposes a similar model in their first example titled “Continuous individual covariates.” An important distinction though is the allowance, for each covariate, to have a different set of coefficients, $\boldsymbol{\beta}_j$. King et al. (2016) justifies a single coefficient per covariate as the model they develop is in the context of animal populations. In animal

populations, each list is a different capturing from the same population but at various time points. In order to assume a closed population, the lists should be collected in a relatively short time period. As a result, the author argues that any time varying effect from the covariates should be limited. On the other hand, we approach this problem in the context of human populations, where lists are not necessarily different time points but instead different data collectors or databases. As a result, we expect individual characteristics (covariates) to have different impacts on the capture probability depending on the list.

In order to complete the Bayesian model, priors must be assigned to unknown parameters N , β_j , and ϕ_h . For N , we use the Jeffrey’s prior (Jeffreys, 1967), $p(N) = \frac{1}{N}$, which conveniently results in a negative binomial distribution for the conditional posterior distribution of N . Other choices of priors typically result in more complicated estimation, especially with data augmentation (see King et al. (2016) for a discussion). For the β_j coefficients, we assign a multivariate normal prior to the set of coefficients for each list with mean of $\mathbf{b} \in \mathcal{R}^{H+1}$ and covariance of $\mathbf{B} \in \mathcal{R}^{(H+1) \times (H+1)}$. Other choices of prior can be used as the problem reduces to Bayesian logistic regression after augmenting the missing data. We select the multivariate normal prior as it is the same prior used in the Bayesian logistic regression Monte Carlo Markov Chain (MCMC) sampling scheme proposed in (Polson et al., 2013), which makes implementation simple. We can be flexible in our choice of prior distribution for the β_j coefficients as long as a suitable method exists for drawing samples from the conditional posterior distribution for β_j . Lastly, for now, assume ϕ_h is known and thus can be treated as a hyperparameter (this will be further addressed in subsection 4).

Plugging in the likelihood distribution as described in Equation 6 and the aforementioned

priors into Equation 4 yields the posterior

$$\begin{aligned}
p(N, \boldsymbol{\beta}, \mathcal{X}_{mis} | \mathcal{Y}_{obs}, \mathcal{X}_{obs}) &\propto \left[\binom{N}{n} \prod_{i=1}^n \prod_{j=1}^J \lambda_{ij}(\mathbf{x}_i)^{y_{ij}} (1 - \lambda_{ij}(\mathbf{x}_i))^{1-y_{ij}} \prod_{i=n+1}^N \prod_{j=1}^J (1 - \lambda_{ij}(\mathbf{x}_i)) \right] \\
&\times \left[\prod_{i=1}^n \mathbf{g}(\mathbf{x}_i | \boldsymbol{\phi}_h) \right] \times \left[\prod_{i=n+1}^N \mathbf{g}(\mathbf{x}_i | \boldsymbol{\phi}_h) \right] \times \left[\frac{1}{N} \right] \\
&\times \left[\prod_{j=1}^J \left(\frac{1}{2\pi} \right)^{H/2} |\mathbf{B}|^{-1/2} e^{-\frac{1}{2}(\mathbf{b}-\boldsymbol{\beta}_j)^T \mathbf{B}^{-1}(\mathbf{b}-\boldsymbol{\beta}_j)} \right]. \tag{9}
\end{aligned}$$

3.2.2 Estimation of the BLRCR

An exact analytic solution for the posterior is intractable, so we implement the MCMC algorithm of Gibbs Sampling, where sequential draws from $\boldsymbol{\beta}$, N , and \mathcal{X}_{mis} are taken conditional on all other parameters.

- 1) Sample $\boldsymbol{\beta}$. For this stage \mathcal{X}_{mis} and \mathcal{X}_{obs} are both known implying N is known as well.

Therefore the sampling equation reduces to

$$p(\boldsymbol{\beta} | \mathcal{Y}, N, \mathcal{X}_{mis}, \mathcal{X}_{obs}) = p(\boldsymbol{\beta} | \mathcal{Y}, \mathcal{X}), \tag{10}$$

which is simply the posterior distribution of Bayesian logistic regression. Posterior samples can be obtained from Equation 10 by first sampling a latent variable from the Polya-Gamma distribution and using this latent variable in the mean and covariance function of a multivariate normal (Polson et al., 2013).

- 2) Sample N and \mathcal{X}_{mis} . Adding up the number of missing covariates, n_0 , with the number of observed covariates, n , fully determines $N = n_0 + n$. This issue makes it impossible to compute the standard Gibbs sampling equation for N . To get around this complication, we sample the parameters simultaneously (Basu, 2001),

$$p(N, \mathcal{X}_{mis} | \mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs}) \propto p(N | \mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs}) p(\mathcal{X}_{mis} | N, \mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs}). \tag{11}$$

From [Equation 11](#), observe the joint distribution of N and \mathcal{X}_{obs} can be decomposed into two parts from which can be sampled. The first part of this equation is the joint distribution of N and \mathcal{X}_{mis} marginalized over the missing covariates. The second part is the distribution of the missing covariates where the number of missing covariates, n_0 , is known.

- i. Sample $N \sim p(N|\mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs})$.

$$\begin{aligned}
p(N|\mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs}) &= \int_{\mathbf{x}_{n+1}} \cdots \int_{\mathbf{x}_N} p(N, \mathcal{X}_{mis}|\mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs}) d\mathbf{x}_{n+1} \cdots d\mathbf{x}_N \\
&\propto \int_{\mathbf{x}_{n+1}} \cdots \int_{\mathbf{x}_N} \left[\binom{N}{n} \prod_{i=n+1}^N \prod_{j=1}^J (1 - \lambda_{ij}) \right] \left[\prod_{i=n+1}^N \mathbf{g}(\mathbf{x}_i|\boldsymbol{\phi}) \right] \left[\frac{1}{N} \right] d\mathbf{x}_{n+1} \cdots d\mathbf{x}_N \\
&= \frac{(N-1)!}{(N-n)!n!} \left[\int_{\mathbf{x}} \mathbf{g}(\mathbf{x}|\boldsymbol{\phi}) \prod_{j=1}^J (1 - \lambda_{ij}) d\mathbf{x} \right]^{N-n} \\
&\propto \binom{N-1}{n-1} \left[\underbrace{E_{\mathbf{g}(\boldsymbol{\theta}_n)} \left[\prod_{j=1}^J (1 - \lambda_{ij}) \right]}_{\equiv \rho} \right]^{N-n} \tag{12}
\end{aligned}$$

Instead of sampling N , sample $n_0 = N - n$.

$$\begin{aligned}
p(n_0|\mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs}) &\propto \binom{n_0 + n - 1}{n - 1} \rho^{n_0} \\
&\propto \binom{n_0 + n - 1}{n - 1} \rho^{n_0} \underbrace{(1 - \rho)^n}_{=constant} \\
&= \text{NegativeBinomial}(n, 1 - \rho). \tag{13}
\end{aligned}$$

The distribution of n_0 follows a negative binomial with parameter n for the number of "successes" and $1 - p$ as the "success" rate. In this context, a "success" is defined as an observation being unobserved. The value of ρ is defined in [Equation 12](#) and can be computed via numerical integration or estimated via a Monte Carlo within MCMC step as in [Bonner and Schofield \(2014\)](#). After

computing ρ and sampling n_0 through Equation 13, find $N = n_0 + n$.

- ii. Sample \mathcal{X}_{mis} . The missing observation, \mathbf{x}_i , is drawn independently of all other covariates, so \mathbf{x}_i does not depend on any other $\mathbf{x}_i \in \mathcal{X}_{obs} \cup \mathcal{X}_{mis}$. Also, by definition, if $\mathbf{x}_i \in \mathcal{X}_{mis}$, then $\mathbf{y}_i = \mathbf{0}$. Therefore, the distribution to be sampled is

$$\begin{aligned} p(\mathbf{x}_i|N, \mathcal{Y}, \boldsymbol{\beta}, \mathcal{X}_{obs}) &= p(\mathbf{x}_i|\mathbf{y}_i = \mathbf{0}, \boldsymbol{\beta}) \\ &\propto \mathbf{g}(\mathbf{x}_i|\boldsymbol{\phi}) \prod_{j=1}^J (1 - \lambda_{ij}(\mathbf{x}_i)) \end{aligned} \quad (14)$$

To sample from Equation 14, we use rejection sampling of a truncated distribution. To do this, first draw a sample \mathbf{x}_i from the distribution of the covariates, $\mathbf{g}(\boldsymbol{\phi})$. Next, accept the sample with probability of \mathbf{x}_i being missing, i.e., $\prod_{j=1}^J (1 - \lambda_{ij}(\mathbf{x}_i))$. If the sample is not accepted, reject it, and draw another sample from $\mathbf{g}(\boldsymbol{\phi})$. Repeat this until we obtain $n_0 = N - n$ missing covariates. On average, this sampling techniques requires us to draw N total covariates for each sample.

3.3 Simulation Analysis

In this subsection, we examine the Bayesian Logistic Regression Capture-Recapture (BLRCR) model and estimation procedure using a variety of simulations and compare it with other popular methodologies. Of interest is the computational efficiency and quality of the model. Since the estimation procedure utilizes a MCMC methodology, the concern is in the mixing of the sampled posterior distribution. Typically, as the model becomes more complex with a greater number of parameters, the need arises for longer MCMC chains in order to converge to the target posterior stationary distribution. Meanwhile, assuming the model can converge, we also assess the quality of the procedure under various measures of accuracy (bias, mean-squared error, capture percentage) and precision (relative length of interval).

While the primary purpose is to examine the results of the Bayesian Logistic Regression Capture-Recapture (**BLRCR**) model, we use four other popular capture-recapture algorithms as a comparison. First, we consider the conditional maximum likelihood logistic regression (**cMLCR**) which we implement using the **VGAM** package in **R** (see subsection 4). The package provides asymptotic estimates for the standard error and suggests, through examples, to construct confidence intervals assuming normality. Instead, utilize the package to get our point estimate and proceed with a semiparametric bootstrap for the confidence intervals (Zwane and van der Heijden, 2003). Second, we consider the ubiquitous hierarchical log-linear (**Log Linear - BIC**) modelling technique (Fienberg, 1972) using the **Rcapture** package in **R**. Keep in mind that this approach does not use covariates, but attempts to model list dependency directly through list interactions. Since the approach is hierarchical with 2^J different model constructions, all are calculated and the one with the lowest BIC is selected. Third, we use the Bayesian Non-Parametric Latent-Class Capture-Recapture (**LCMCR**) algorithm in Manrique-Vallier (2016) which is another technique that does not allow covariates but uses a Bayesian nonparametric approach to account for unobserved heterogeneity (see Section 11: Appendix B for a summary of this approach). The fourth technique is a simple independence model (**Independent**) where it is assumed there is no list or individual heterogeneity. Conveniently, when the number of latent classes is set equal to one, the LCMCR model collapses into an independence model, effectively giving a Bayesian independence sampler. The BLRCR model is designed to be flexible and extensible (see sections 4, 5, and 6), hence there are several choices in its implementation. While we will implement these various extensions in upcoming sections, for now, we will assume the true distribution of the covariates is known. This does not, however, avoid the necessity of specifying hyperparameters for the Bayesian models. In order to be consistent throughout the entirety of this dissertation, we use the same relatively uninformative prior specifications throughout all sections. Table 1 lists these hyperparameter specifications for both the

BLRCR and LCMCR models.

	Method	
Hyperparameter	SP-BLRCR	LCMCR
a	0.25	0.25
b	0.25	0.25
ν_0	3	
κ_0	1	
$\boldsymbol{\mu}_0$	(0,0)	
$\boldsymbol{\Lambda}_0$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$	

Table 1: Hyperparameter specifications for the SP-BLRCR and LCMCR methodologies.

Since the BLRCR model uses logistic regression with a normal distribution prior on the coefficients, this also requires a prior for the mean and covariance, which are set to

$$\mathbf{b} = \mathbf{0} \text{ and } \mathbf{B} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}.$$

For each set of simulations, a table with results can be found. For each algorithm a point estimate of the population size, N , is computed. We have some options for what to use as the point estimate, \hat{N} , for the MCMC algorithms, but we elect to use median of the sampled posterior. The column *Mean N* records the average \hat{N} across the simulations and the column *Mean N%* takes it as a percentage of N . Hence, if the true population size is $N = 2000$, a *Mean N* = 2000 and a *Mean N%* = 1.000 would indicate an unbiased estimate. We also consider the accuracy in two other ways: the column *MSE* records the estimated mean squared error of the point estimate and the column *CI%* records the percentage of times the 95% interval captured the true population number, N . In order to get a measure of interval's precision, we compute the average of the simulated 95% confidence/credible interval widths as a percentage of N , which is recorded in column

CI Width. Ideally, we would want a model with high precision, i.e. a small interval width; but still maintains the ability to be highly accurate.

3.3.1 Simulations with Varying Levels of List Dependency

The objective of this subsection is to evaluate the BLRCR and other models under varying coefficients that induce different levels of list dependency. All of this list dependency is really just heterogeneity in the capture probabilities that can be fully explained by the covariates. The models that do not use covariates (Log Linear, LCMCR, and Independent) will not or will struggle to detect the heterogeneity leading to biased results. On the other hand, the BLRCR and cMLCR will perform well as they incorporate the covariates into their modelling.

We simulate data in accordance with [Equation 6](#) using two covariates ($H=2$) and four lists ($J = 4$). The covariates are drawn from two independent standard normal distributions. Three sets of β coefficients are chosen that create “negative”, “moderate”, and “positive” dependency between the lists and are described in [Table 2](#). To clarify, all of the lists regardless of the coefficients are conditionally independent given the covariates. We simulate an additional dataset with all of the slope coefficients set to 0, thereby creating lists with independent capture probabilities.

"Negative"				"Moderate"				"Positive"			
List (j)	β_{0j}	β_{1j}	β_{2j}	List (j)	β_{0j}	β_{1j}	β_{2j}	List (j)	β_{0j}	β_{1j}	β_{2j}
1	-2	-1	1	1	-2	-1	1	1	-2	-1	1
2	-2	1	-1	2	-2	1	-1	2	-2	-1	1
3	-2	1	1	3	-2	-1	1	3	-2	-1	1
4	-2	-1	-1	4	-2	1	-1	4	-2	-1	1

Table 2: Coefficients for Simulated Data

β	Method	N%	MSE	CI Width	CI %
"Moderate"	BLRCR($K = 1, H_\omega = 1$)	0.983	82.5	0.156	94.0
	BLRCR($K = 20, H_\omega = 1$)	0.988	82.6	0.162	95.0
	BLRCR($K = 20, H_\omega = 20$)	0.997	84.9	0.204	99.0
	cMLCR	1.002	87.1	0.163	92.0
	Log Linear (BIC)	1.399	1669.3	0.458	0.0
	LCMCR	0.851	303.4	0.148	11.0
	Independent	1.035	103.1	0.157	90.0
"Negative"	BLRCR($K = 1, H_\omega = 1$)	0.987	79.0	0.149	97.0
	BLRCR($K = 20, H_\omega = 1$)	0.990	76.3	0.151	100.0
	BLRCR($K = 20, H_\omega = 20$)	1.006	80.5	0.189	99.0
	cMLCR	1.000	77.1	0.150	99.0
	Log Linear (BIC)	0.999	81.0	0.159	98.0
	LCMCR	1.008	99.3	0.284	98.0
	Independent	1.208	426.7	0.206	0.0
"Positive"	BLRCR($K = 1, H_\omega = 1$)	0.942	188.0	0.350	91.0
	BLRCR($K = 20, H_\omega = 1$)	0.952	168.1	0.354	93.0
	BLRCR($K = 20, H_\omega = 20$)	0.978	158.0	0.385	97.0
	cMLCR	1.025	207.3	0.389	94.0
	Log Linear (BIC)	0.899	301.4	0.322	77.0
	LCMCR	0.699	608.1	0.236	4.0
	Independent	0.566	868.4	0.048	0.0
"Independent"	BLRCR($K = 1, H_\omega = 1$)	0.981	120.2	0.247	96.0
	BLRCR($K = 20, H_\omega = 1$)	0.980	119.2	0.247	96.0
	BLRCR($K = 20, H_\omega = 20$)	1.009	128.6	0.370	100.0
	cMLCR	1.011	126.6	0.261	97.0
	Log Linear (BIC)	1.002	121.5	0.256	98.0
	LCMCR	0.993	115.7	0.260	99.0
	Independent	0.989	115.9	0.247	97.0

Table 3: Results of 100 capture-recapture simulations per coefficient set using $N = 2000$ with "Moderate", "Negative", "Positive", and "Independent" list dependency with two standard normal covariates.

Table 3 shows the results from 100 simulated datasets for each set of coefficients, β . While all of the models performed well when the slope coefficients were set to 0, i.e. "Independent", the methods using covariates (BLRCR and cMLRCR) showed substantially less bias than the other methods when coefficients impacted the capture probability, i.e. "Moderate", "Negative", and "Positive". This illustrates the importance of including covariates in the estimation process. Nevertheless, the two methods that account

for unknown heterogeneity, log linear and LCMCR, still performed remarkably well when the list dependency was set to be “negative.” Unfortunately, “positively” induced list dependency led to a substantial decline in performance.

3.3.2 Simulations with Different Sized Populations

In this subsection, we examine the effects of differing population sizes on the estimations with two objectives to evaluate in mind: performance and consistency. For the most part, we will see that all models will become less biased as the population size increases (except assuming independence). We will also see that using the BLRCR will tend to have a lower MSE than the cMLCR approach, especially with smaller population sizes.

Data was simulated with varying population sizes, $N = [200, 500, 1000, 2000, 5000, 10000]$, with the results summarized in [Table 4](#). For consistency, all datasets were constructed using the coefficients yielding “moderate” dependency and with two standard normal covariates. Before examining the results, it should be pointed out that there were considerable issues with using the log linear approach with the smaller sample sizes. While a point estimate was always able to be obtained, often times, the `Rcapture` package would simply report a lower bound on the upper limit of the confidence interval. Instead of making a decision on whether to use that bound as the upper limit, we simply report the interval width as not available (NA).

With this consideration in mind, we highlight a few interesting observations. Most importantly, the BLRCR and cMLCR methods tended to outperform the other methods in terms of mean squared error (MSE), 95 % credible interval width (CI width), and the 95% credible interval capture percentage (CI %). There were a handful of occasions where that was not the case but appears to be attributed to the choice of simulation parameters. Notably, the independence model outperformed all other models when $N = 1000$, but as the population grew, it became substantially worse and is actually the worst option when

$N = 10000$. Additional simulations (not shown) performed with different sets of coefficients revealed the independence model to be a substantially worse choice than BLRCR regardless of population size.

A comparison of the three BLRCR methods and cMLCR reveal relatively similar performances across population sizes. Keep in mind, the simulation used in [Table 3](#) simply used a single multivariate normal with no heterogeneity. The additional noise of assuming an infinite mixture of normal distributions and/or heterogeneity in capture probabilities seems to have had little to no impact on the estimation. [Figure 1](#) shows a correlation plot between the 100 estimates for N using the three BLRCR methods and cMLCR with the true population size set to $N = 2000$. Notice, the estimates are highly correlated, and BLRCR tends to return smaller estimated values than cMLCR. An important result of [Alho \(1990\)](#), is that the cMLCR estimator is consistent but may be biased with small sample sizes. From the simulations, it certainly appears this is true for $N=200$, but the level of bias quickly disappears with $N=500$ and above. Similarly, the BLRCR seems to share this quality of consistency but the bias is greater albeit the opposite direction. Nevertheless, as we would expect with Bayesian methods, we are trading some bias for a reduction in variance. The mean squared error (MSE) tends to be smaller than the other methods.

N	Method	N%	MSE	CI Width	CI %
200	BLRCR($K = 1, H_\omega = 1$)	0.887	28.8	0.374	78.0
	BLRCR($K = 20, H_\omega = 1$)	0.883	29.4	0.375	78.0
	BLRCR($K = 20, H_\omega = 20$)	0.914	26.4	0.459	92.0
	cMLCR	1.057	36.7	0.635	96.0
	Log Linear (BIC)	2.316	472.6	NA	23.0
	LCMCR	0.854	35.7	0.434	74.0
	Independent	1.004	22.0	0.481	97.0
500	BLRCR($K = 1, H_\omega = 1$)	0.940	44.0	0.271	91.0
	BLRCR($K = 20, H_\omega = 20$)	0.943	42.7	0.280	90.0
	BLRCR($K = 20, H_\omega = 20$)	0.966	37.6	0.342	96.0
	cMLCR	1.010	41.0	0.340	95.0
	Log Linear (BIC)	2.084	690.8	1.826	1.0
	LCMCR	0.840	84.1	0.253	43.0
	Independent	1.017	35.4	0.306	96.0
1000	BLRCR($K = 1, H_\omega = 1$)	0.972	58.3	0.210	93.0
	BLRCR($K = 20, H_\omega = 20$)	0.976	56.9	0.217	94.0
	BLRCR($K = 20, H_\omega = 20$)	0.991	57.9	0.267	97.0
	cMLCR	1.008	60.1	0.235	95.0
	Log Linear (BIC)	1.808	1106.6	0.957	0.0
	LCMCR	0.852	154.5	0.190	20.0
	Independent	1.033	62.8	0.222	92.0
2000	BLRCR($K = 1, H_\omega = 1$)	0.983	82.5	0.156	94.0
	BLRCR($K = 20, H_\omega = 1$)	0.988	82.6	0.162	95.0
	BLRCR($K = 20, H_\omega = 20$)	0.997	84.9	0.204	99.0
	cMLCR	1.002	87.1	0.163	92.0
	Log Linear (BIC)	1.399	1669.3	0.458	0.0
	LCMCR	0.851	303.4	0.148	11.0
	Independent	1.035	103.1	0.157	90.0
5000	BLRCR($K = 1, H_\omega = 1$)	0.993	123.7	0.107	96.0
	BLRCR($K = 20, H_\omega = 1$)	0.997	120.2	0.112	98.0
	BLRCR($K = 20, H_\omega = 20$)	1.002	127.2	0.145	100.0
	cMLCR	1.001	122.1	0.102	96.0
	Log Linear (BIC)	1.024	1651.3	0.326	62.0
	LCMCR	0.898	763.4	0.277	37.0
	Independent	1.039	224.4	0.100	65.0
10000	BLRCR($K = 1, H_\omega = 1$)	0.996	174.0	0.084	97.0
	BLRCR($K = 20, H_\omega = 1$)	0.998	177.0	0.087	98.0
	BLRCR($K = 20, H_\omega = 20$)	0.997	197.0	0.109	99.0
	cMLCR	0.999	178.9	0.072	94.0
	Log Linear (BIC)	1.000	801.5	0.278	95.0
	LCMCR	0.957	1491.2	0.385	64.0
	Independent	1.039	420.8	0.071	35.0

Table 4: Results of 100 capture-recapture simulations per varying population sizes using "Moderate" list dependency and two standard normal covariates.




Figure 1: Correlation Plot of the 100 estimates for N using BLRCR and cMLCR when simulating from two independent standard normal distributions with moderate dependency between lists and a true population size of $N=2000$.

Overall, the BLRCR model is shown to be consistent. It also tends to outperform or equally perform the other methods in terms of MSE, CI width, and CI%.

4 Selecting a Distribution for the Covariates

When the probability that an individual is captured at least once is dependent upon the covariates, the observed covariate distribution will differ from population covariate distribution. Simply using logistic regression on the observed data would lead to biased coefficients, which in turn would lead to bias in the population estimation. The algorithm

presented in subsection 3.2 alleviates this problem by concatenating samples of the missing covariates with the observed data before estimating the coefficients. Unfortunately, this presents an additional burden on the user of specifying a distribution for the missing covariates, $\mathbf{g}(\phi)$.

4.1 Specifying a Distribution for the Covariates

If the distribution is known including parameters, then we can simply use the aforementioned algorithm. If the distribution is known except for the parameters, then we could specify that distribution along with priors on the parameters. For example, Royle (2009) uses a single covariate and specifies a normal distribution with a normal distribution prior for the mean and a gamma prior for the inverse variance. One could take a multivariate version of this approach by specifying,

$$\mathbf{x}_i \stackrel{iid}{\sim} \text{MVNormal}(\boldsymbol{\mu}, \boldsymbol{\Sigma}), \quad (15)$$

with conjugate priors,

$$\boldsymbol{\Sigma} \sim \text{InvWishart}(\nu_0, \boldsymbol{\Lambda}_0^{-1}) \quad (16)$$

$$\boldsymbol{\mu} | \boldsymbol{\Sigma} \sim \text{MVNormal}(\boldsymbol{\mu}_0, \boldsymbol{\Sigma} / \kappa_0) \quad (17)$$

This would add two additional sampling stages to the algorithm (see Gelman et al. (2014)):

- 1) Sample $\boldsymbol{\Sigma}$. Define the sufficient statistics, $\bar{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i$ and

$\mathbf{S} = \sum_{i=1}^N (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T$. Then,

$$\boldsymbol{\Sigma} \sim \text{InvWishart}(\nu_N, \boldsymbol{\Lambda}_N^{-1}), \quad (18)$$

where $\nu_N = \nu_0 + N$ and $\boldsymbol{\Lambda}_N = \boldsymbol{\Lambda}_0 + \boldsymbol{S} + \frac{\kappa_0 N}{\kappa_0 + N}(\bar{\boldsymbol{x}} - \boldsymbol{\mu}_0)(\bar{\boldsymbol{x}} - \boldsymbol{\mu}_0)^T$.

2) Sample $\boldsymbol{\mu}_k | \boldsymbol{\Sigma}_k$ for $k = 1, \dots, K^*$. Using the same defined terms in the previous step,

$$\boldsymbol{\mu} | \boldsymbol{\Sigma} \sim \text{MVNormal}(\boldsymbol{\mu}_N, \boldsymbol{\Sigma} / \kappa_N), \quad (19)$$

where $\boldsymbol{\mu}_N = \frac{\kappa_0}{\kappa_0 + N} \boldsymbol{\mu}_0 + \frac{N}{\kappa_0 + N} \bar{\boldsymbol{x}}$ and $\kappa_N = \kappa_0 + N$.

Different specified distributions would require different sampling procedures. In a supplementary document, [Royle \(2009\)](#) considers other types of covariate distributions and found some variations in the inference. Unfortunately, knowing what distribution to specify can be difficult as rarely would we know this distribution. Deciding on a proposal distribution based upon the observed distribution can be also dangerous or misleading as the observed distribution is a truncation of the true distribution. Not only are the missing covariates not missing at random but we do not even know how many are missing. As we will shown in subsection 4.4, misspecifying the distribution can lead to inaccurate estimations.

4.2 Conditional Likelihood

One approach to the problem is estimating the coefficients using the conditional maximum likelihood, where the likelihood function to be maximized is conditioned on the probability an individual is observed at least once ([Alho, 1990](#); [Huggins, 1989](#)). This avoids the necessity of specifying a distribution for the covariates but can lead to unstable results when the observed distribution of the covariates is dissimilar to the population distribution of the covariates ([Tilling and Sterne, 1999](#)). An article by [Yee et al. \(2015\)](#) shows an easy

way to implement the technique using the **VGAM** package in R. The technique works in two stages. First, use generalized linear models with a positive Bernoulli family to estimate the coefficients. This allows us to obtain fitted values for the probability that each individual in the dataset is missing. Second, use those fitted values in the Horvitz-Thompson estimator ([Horvitz and Thompson, 1952](#)) to estimate the population size, N . One could take a Bayesian approach to the logistic regression and assign prior distributions to the parameters and estimate the population in a similar way. We derive equations for finding the maximum a posteriori (MAP) estimate in [Section 10: Appendix A](#) using gradient ascent.

4.3 Nonparametric Distributions

An alternative approach to the methods described above is to specify a nonparametric distribution for the covariates. Using a nonparametric method has the added consequence of increasing the complexity of the model which may lead to additional difficulties in the mixing of the posterior. Nevertheless, it comes with multiple benefits in allowing us to maintain use of a full likelihood approach, continued extensibility, and (most importantly) flexibility in the fitting of the covariates. In this subsection, we present two methods for modeling the covariates in a non-parametric fashion: The Dirichlet Process Mixture of Normal Distributions and the Bayesian Bootstrap.

4.3.1 Dirichlet Process Mixture of Normal Distributions

The first method is to fit a Dirichlet process mixture of normal distributions. This nonparametric approach fits a potentially infinite number of multivariate normal distributions to the data. The following is a description of the methodology and is adapted from [Gelman et al. \(2014\)](#). This model introduces a latent variable, z , which is a latent parameter determining the mean and covariance matrix from which the observed covariate,

\mathbf{x}_i , is drawn. The generative process for each observation's covariate can then be summarized as,

$$\mathbf{x}_i|z_i \stackrel{iid}{\sim} \text{MVNormal}(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \text{ for } i = 1, \dots, N \quad (20)$$

$$z_i \stackrel{iid}{\sim} \text{Discrete}(\{1, 2, \dots\}, (\pi_1, \pi_2, \dots)) \text{ for } i = 1, \dots, N. \quad (21)$$

Therefore, the probability density function of the covariates can be formally written as,

$$g(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = \prod_{i=1}^N \prod_{k=1}^{\infty} \pi_k \text{MVNormal}(\mathbf{x}_i|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k). \quad (22)$$

In order to have a fully Bayesian approach, we assign priors to the unknown parameters.

$$\boldsymbol{\Sigma}_k \sim \text{InvWishart}(\nu_0, \boldsymbol{\Lambda}_0^{-1}) \quad (23)$$

$$\boldsymbol{\mu}_k|\boldsymbol{\Sigma}_k \sim \text{MVNormal}(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_k/\kappa_0) \quad (24)$$

$$(\pi_1, \pi_2, \dots) \sim \text{SB}(\alpha) \quad (25)$$

$$\alpha \sim \text{Gamma}(a_\alpha, b_\beta), \quad (26)$$

where $\text{SB}(\alpha)$ is the stick breaking process, ([Ishwaran and James, 2001](#)). While we model the number of latent classes as infinite, we can approximate this infinite mixture problem by setting a sufficiently large upperbound to the number of latent classes, K^* , and solving this finite-dimensional problem. The stick breaking prior can therefore be defined as $\pi_k = V_k \prod_{l < k} (1 - V_l)$, where $V_1, \dots, V_{K^*-1} \sim \text{Beta}(1, \alpha)$ and $V_{K^*} = 1$. This upper bound, K^* , should not be thought of as a parameter as it should have no impact on the estimation process as long as the value is set sufficiently large enough.

Using this generative scheme adds five additional sampling stages to the algorithm presented in subsection 3.2.

- 1) Sample z_i for $i = 1, \dots, N$. The latent class label takes integer values from $1, \dots, K^*$.

To compute the probability of each latent class label for each i ,

$$P(z_i = k) = \frac{\pi_k \text{MVNormal}(\mathbf{x}_i | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{l=1}^{K^*} \pi_l \text{MVNormal}(\mathbf{x}_i | \boldsymbol{\mu}_l, \boldsymbol{\Sigma}_l)}. \quad (27)$$

- 2) Sample $\boldsymbol{\Sigma}_k$ for $k = 1, \dots, K^*$. Define $N_k = \sum_{i=1}^N 1_{z_i=k}$ which is a count of the number of individuals in the population belonging to latent class, k . Also, define the sufficient statistics, $\bar{\mathbf{x}}_k = \frac{1}{N_k} \sum_{i=1}^{N_k} \mathbf{x}_{ik}$ and $\mathbf{S}_k = \sum_{i=1}^{N_k} (\mathbf{x}_{ik} - \bar{\mathbf{x}}_k)(\mathbf{x}_{ik} - \bar{\mathbf{x}}_k)^T$. Then,

$$\boldsymbol{\Sigma}_k \sim \text{InvWishart}(\nu_{N_k}, \boldsymbol{\Lambda}_{N_k}^{-1}), \quad (28)$$

where $\nu_{N_k} = \nu_0 + N_k$ and $\boldsymbol{\Lambda}_{N_k} = \boldsymbol{\Lambda}_0 + \mathbf{S}_k + \frac{\kappa_0 N_k}{\kappa_0 + N_k} (\bar{\mathbf{x}}_k - \boldsymbol{\mu}_0)(\bar{\mathbf{x}}_k - \boldsymbol{\mu}_0)^T$.

- 3) Sample $\boldsymbol{\mu}_k | \boldsymbol{\Sigma}_k$ for $k = 1, \dots, K^*$. Using the same defined terms in the previous step,

$$\boldsymbol{\mu}_k | \boldsymbol{\Sigma}_k \sim \text{MVNormal}(\boldsymbol{\mu}_{N_k}, \boldsymbol{\Sigma}_k / \kappa_{N_k}), \quad (29)$$

where $\boldsymbol{\mu}_{N_k} = \frac{\kappa_0}{\kappa_0 + N_k} \boldsymbol{\mu}_0 + \frac{N_k}{\kappa_0 + N_k} \bar{\mathbf{x}}_k$ and $\kappa_{N_k} = \kappa_0 + N_k$.

- 4) Sample $(\pi_1, \pi_2, \dots, \pi_{K^*})$ for $k = 1, \dots, K^*$. Begin by drawing a sample from each of the stochastic components,

$$V_k \sim \text{Beta} \left(1 + N_k, \alpha + \sum_{l=k+1}^K N_l \right) \text{ for } k = 1, \dots, K^* - 1. \quad (30)$$

Set $V_{K^*} = 1$. Then, $\pi_k = V_k \prod_{l < k} (1 - V_l)$ for all $k = 1, \dots, K^*$.

- 5) Sample α .

$$\alpha \sim \text{Gamma} \left(a_\alpha + K^* - 1, b_\alpha - \sum_k^{K^*-1} \ln(1 - V_k) \right). \quad (31)$$

It should be noted that the normal distribution specified in subsection 4.1 is a special case of this specification where $K^* = 1$. We will show in section 4.4 that the assumption of a single normal may lead to bias in the estimation process when the actual distribution is not normal. Using an infinite mixture of normal distributions tends to perform better, but still struggles when the distribution is discrete or is far from normally distributed.

4.3.2 The Bayesian Bootstrap

An alternative approach would be to apply a discrete distribution with support at the observed, distinct values, \mathbf{d}_k , of the covariates,

$$P(\mathbf{x}_i = \mathbf{d}_k) = \psi_k, \quad \sum_{k=1}^K \psi_k = 1. \quad (32)$$

We attach a Dirichlet prior to the probabilities, ψ_k , such that

$$(\psi_1, \dots, \psi_K) \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_K) \quad (33)$$

The posterior distribution for the ψ_k is then,

$$(\psi_1, \dots, \psi_K | \mathcal{X}) \sim \text{Dirichlet}(\alpha_1 + \sum_{i=1}^N I_{d_1}(\mathbf{x}_i), \dots, \alpha_K + \sum_{i=1}^N I_{d_K}(\mathbf{x}_i)) \quad (34)$$

where $I_{d_k}(\mathbf{x}_i) = 1$ is an indicator function for $\mathbf{x}_i = \mathbf{d}_k$ and $I_{d_k}(\mathbf{x}_i) = 0$ otherwise.

Although we are not performing an actual bootstrapping procedure, we refer to this modelling setup as the Bayesian Bootstrap (Rubin, 1981) for a couple of reasons. First, the construction is quite similar to the theory behind the Bayesian Bootstrap in terms of

placing a Dirichlet distribution on the observed values in the sample. Second, this methodology samples values for the missing observations in the same spirit as bootstrapping. Using the n observed covariates, we sample, with replacement, n_0 values from a discrete distribution based on the observed distribution and use these values in place of the missing covariates.

To implement this model construction for the covariates, we only need to add one additional sampling stage to our estimation procedure. Noting that we use all covariates, not just the observed, we draw the probabilities, (ψ_1, \dots, ψ_K) in accordance to equation 34. Once the probabilities, (ψ_1, \dots, ψ_K) , are drawn, the covariate distribution is fully specified, so the rejection sampling method as displayed in equation 14 can be performed.

A common approximation to the Bayesian Bootstrap is to sample directly from the empirical distribution. When $(\alpha_1, \dots, \alpha_K) = \alpha = 0$, one could skip the Dirichlet distribution sampling stage and simply draw, with replacement, from the current complete covariate data, \mathcal{X} . While this is a good approximation, it will have less smoothing than its Bayesian counterpart.

This Bayesian Bootstrap method addresses some of the potential issues in the covariate specification, but it also introduces new issues as well. This method provides an “all-purpose” solution in that it can be used when the covariate distribution is discrete, continuous, or both. Nevertheless, while it alleviates some of the issue with data type, no method is truly “all-purpose” as any method will provide some sort of constraint. The issues this particular model tends to have aligns closely with those of the Bayesian Bootstrap. For example, the missing values that are sampled only has support at values in the observed empirical distribution. What happens if some of the potential values that are missing do not appear in the observed distribution? This is a concern for when the covariate distribution is discrete but a probabilistic certainty when the distribution is continuous. This should be a concern but as Rubin (1981) states “serious data analysis

should always include serious consideration of model constraints” and that no method, even the bootstrap, is without constraints. Furthermore, we note that the end goal is to estimate the size of the population, so the exact modelling of the covariate distribution is not as important as long as it does not bias the estimation. Using the rejection sampler based on equation 14 combined with $(\alpha_1, \dots, \alpha_K) = \alpha = 0$, the probability of each draw is appropriately weighted based on the probability of going unobserved. This has similar intuition as the mechanic behind the Horvitz-Thompson estimator in the conditional likelihood approach where each covariate essentially gets magnified based on its estimated probability of being missing.

4.4 Simulations using Different Covariate Distributions

EXTRA STUFF BELOW WRITTEN HERE, MAY NEED TO BE MOVED By default, we do not use any regularization (see subsection 6) unless specifically stated. For the covariate distribution, there are three possibilities including the mixture of multivariate normal distributions (**mixture**), the Bayesian Bootstrap (**bootstrap**), and the approximation of the Bayesian Bootstrap that simple samples from the empirical distribution (**empirical**). In general, the maximum number of mixture components allowed by the mixture of multivariate normal distributions (K^*) will be specified but is intended to be set intentionally large enough so that it does not influence the estimation. The exception would be the case where we desire a single multivariate normal distribution ($K^* = 1$). Similarly, when latent intercepts are used (see subsection 5), the number of latent intercepts (H_ω) is specified.

EXTRA STUFF ABOVE WRITTEN HERE, MAY NEED TO BE MOVED

In this subsection, we examine the impact of different covariate distributions and how they impact the BLRCR model’s estimation. Recall, the BLRCR requires the specification of a covariate distribution, whereas the cMLCR does not. Hence, as we will see below, a

misspecification of the covariate distribution can lead to bias in the inference. This bias will be reduced when using a more flexible covariate distribution like the infinite mixture of normal distributions ($K=20$).

Table 5 presents the results of 100 simulations for three different sets of covariate distributions. The first set of simulations uses the two standard normal distributions seen in the previous subsections. While the cMLCR algorithm is unbiased, all three specifications of the BLRCR algorithm yield a lower MSE. Since the distribution of the covariates is actually normal, it is not surprising that the BLRCR performs well with these simulations. On the other hand, we would expect the second set of covariates, two independent chi-square(1) distributions, to be particularly challenging. The chi-square distribution only has probability mass for nonnegative values, which creates an abrupt cutoff at 0. The third set of covariates includes two different Gamma distributions, Gamma(1,1) and Gamma(3,1). This set of covariates cuts off abruptly on one axis at 0, but not the other. We can therefore think of the three sets of covariates as “normal”, “not normal”, and “near/approximately normal”, respectively.


Distribution	Method	N%	MSE	CI Width	CI %
Normal(0,1)	BLRCR($K = 1, H_\omega = 1$)	0.983	82.5	0.156	94.0
	BLRCR($K = 20, H_\omega = 1$)	0.988	82.6	0.162	95.0
	BLRCR($K = 20, H_\omega = 20$)	0.997	84.9	0.204	99.0
	cMLCR	1.002	87.1	0.163	92.0
	Log Linear (BIC)	1.399	1669.3	0.458	0.0
	LCMCR	0.851	303.4	0.148	11.0
	Independent	1.035	103.1	0.157	90.0
Chi-Square(1)	BLRCR($K = 1, H_\omega = 1$)	1.656	1383.4	0.950	0.0
	BLRCR($K = 20, H_\omega = 1$)	0.967	117.6	0.178	89.0
	BLRCR($K = 20, H_\omega = 20$)	0.986	148.5	0.241	90.0
	cMLCR	1.008	104.8	0.183	92.0
	Log Linear (BIC)	0.892	973.3	0.217	6.0
	LCMCR	0.774	467.0	0.229	31.0
	Independent	0.924	162.8	0.118	32.0
Gamma(1,1)	BLRCR($K = 1, H_\omega = 1$)	1.020	90.7	0.156	94.0
	BLRCR($K = 20, H_\omega = 1$)	0.986	63.4	0.132	99.0
	BLRCR($K = 20, H_\omega = 20$)	0.996	69.6	0.170	100.0
	cMLCR	0.999	55.3	0.123	96.0
	Log Linear (BIC)	1.247	1443.3	0.409	0.0
	LCMCR	0.852	299.3	0.106	7.0
	Independent	0.863	276.6	0.054	0.0

Table 5: Results of capture-recapture algorithms with simulations using different covariate distributions with "moderate" dependency between lists and $N=2000$.


As we saw in the previous section, the BLRCR algorithm handles the normally distributed covariates well. For the chi-square distributed covariates, using a single multivariate normal for the covariates results in poor performance with with an approximate bias of 1.656% of N when $N = 2000$. Even though the covariates are not normally distributed, modeling them as a mixture of normal distributions results in a substantial reduction in the bias for N . [Figure 2](#) illustrates why this may be the case through a partial plotting of the simulated covariates. In all plots, the black dots represent the individuals that are captured at least once in a list. The blue dots represent the individuals that are missing. In the first and second plot in each row, the individuals that are missing are simulated using the BLRCR($K = 1, H_\omega = 1$) and BLRCR($K = 20, H_\omega = 1$) algorithms, respectively. The third panel shows the true missing individuals that are unknown to the algorithm.

With the chi-square distribution, there is no good way to fit a single normal variable that well represents the space. As a result, many missing covariates are populated into low probability density areas resulting in an overestimate of the missing covariates. On the other hand, the BLRCR($K = 20$, $H_\omega = 1$) with it's less rigid covariate assumption, populates the space much better.

Normal(0,1) and Normal(0,1) Covariate Distributions



Chi-Square(1) and Chi-Square(1) Covariate Distributions



Gamma(3,1) and Gamma(1,1) Covariate Distributions




Figure 2: Posterior Distribution of X when simulating various covariate distributions with moderate dependency between lists and $N=2000$.

The parameters of the covariate distribution are sampled based on augmented covariates, not just the observed covariates. This results in uncertainty regarding the ability for the algorithm to correctly identify the covariate distribution. To put the algorithm to the test, we simulate an example with a mixture of three multivariate normal distributions. We use

a population of size 2000 and the “moderate” coefficients. The results of the 100 simulations can be found in [Table 6](#). Despite on average nearly 50% of observations being missing, the algorithm when allowing for up to $K = 20$ mixture normal distributions performs well. Of course, the cMLCR outperforms the algorithm in terms of MSE, but the BLRCR (without also trying to account for heterogeneity) is more precise with a smaller average credible interval width.

Distribution	Method	N%	MSE	CI Width	CI %
Mixture Normal	BLRCR($K = 1, H_\omega = 1$)	0.931	160.1	0.163	67.0
	BLRCR($K = 20, H_\omega = 1$)	0.977	97.9	0.183	96.0
	BLRCR($K = 20, H_\omega = 20$)	0.989	96.3	0.228	99.0
	cMLCR	1.003	90.3	0.191	96.0
	Log Linear (BIC)	1.331	1331.4	0.427	0.0
	LCMCR	0.851	305.5	0.218	25.0
	Independent	1.007	80.0	0.173	96.0

Table 6: Results of capture-recapture algorithms with simulations using different covariate distributions with "moderate" dependency between lists.

A look at the simulated posterior of the covariates in [Figure 3](#) shows once again the benefits of using the mixture of normal distributions. Using a single multivariate normal results in an imputation of covariates in low probability mass spaces, especially between the lower two mixtures. When a mixture of normal distributions is used, the imputation of the covariates is fairly consistent with the true distribution.




Figure 3: Posterior Distribution of the missing X values using algorithm SP-BLRCR.

From this subsection, we saw that depending on the level of the covariate distribution misspecification, the BLRCR model may perform poorly. However, using a non-parametric distribution like the mixture of normal distributions considerably reduced the bias.

5 Conditional Independence and Unobserved Heterogeneity

The model construction of subsection 3.2 assumes conditional independence based on the covariates. In other words, given the information provided by the covariates, the probability of capture on one list is unaffected by another list. If the assumption does not hold, it may lead to biased parameter estimates. One reason the conditional independence may be violated is because the capture probabilities upon a list are directly related to the probability of being on another list, i.e., the assumption of list dependency. This may occur if, for example, one list uses another list as a reference or data is shared between various documentation projects (Manrique-Vallier et al., 2020). This would lead to heavy positive dependence between these two lists. While this is a serious issue, we assume the lists used in the analysis are collected independently.

A second reason conditional independence may be violated is that underlying heterogeneity exists within the population that is not fully accounted for by covariates. Recall, the example cited in subsection 2.1 of individual heterogeneity masquerading as list dependency in the study analyzing extrajudicial killings during the Guatemalan Civil War (Ball et al., 2000). Researchers found that people who were part of Catholic religious communities were more likely to trust Catholic researchers with their stories than with NGO researchers associated with the political left. Similarly, people located in areas associated with the rebel groups were more likely to do the opposite. If this trait is unobserved and not taken into account, it will result in biased β_j coefficients. Further, because of the biased coefficients, the probabilities of capture and the estimate for the population size will be biased as well.

5.1 Modelling the Unobserved Heterogeneity with Latent Classes

We view the problem of unobserved heterogeneity through the lens of a missing covariates problem. There are multiple ways one can induce additional heterogeneity into the modelling. For example, King and Brooks (2008) uses individual random effects following a normal distribution, while Manrique-Vallier (2016) induces heterogeneity by assigning different capture probabilities according to latent group membership. In this subsection, we model hidden heterogeneity also through the use of latent group membership, but instead of each group being assigned a different probability, each group is instead given a different value for an intercept in the linear equation.

Define an indicator vector, ω_i , indicating membership to one of potentially infinite latent groups. As we saw in subsection 4.1 with the infinite mixture of normal distributions, we again make use of the stick-breaking procedure in order to avoid self-selecting the number of the classes to uses. Hence,

$$\phi_\omega \sim \text{SB}(\alpha_\omega) \tag{35}$$

$$\alpha_\omega \sim \text{Gamma}(a_\omega, b_\omega). \quad (36)$$

An additional intercept term is added to the linear equation with the value dependent on the individual's latent class membership. As such, equation 8 then becomes,

$$\lambda_{ij} = \sigma(\mathbf{x}_i^T \boldsymbol{\beta}_j + \boldsymbol{\omega}_i^T \boldsymbol{\beta}_\omega) = \frac{1}{1 + e^{-(\mathbf{x}_i^T \boldsymbol{\beta}_j + \boldsymbol{\omega}_i^T \boldsymbol{\beta}_\omega)}}. \quad (37)$$

This construction leads to four new types of parameters to sample: $\boldsymbol{\omega}_i$, α_ω , $\boldsymbol{\phi}_\omega$, and $\boldsymbol{\beta}_\omega$.

5.2 Updating the Estimation

Since we have four new parameters: $\boldsymbol{\omega}_i$, α_ω , $\boldsymbol{\phi}_\omega$, and $\boldsymbol{\beta}_\omega$; we might expect to need just four new sampling methods. Unfortunately, complications once again arise from the missing covariates of the unobserved individuals. If the individual is observed, the latent group membership, $\boldsymbol{\omega}_i$, can be sampled with corresponding discrete probability,

$$p(\boldsymbol{\omega}_i | \boldsymbol{\phi}_\omega, N, \mathcal{Y}, \boldsymbol{\beta}, \mathbf{X}) \propto \phi_\omega \prod_{j=1}^J \lambda_{ij}^{y_{ij}} (1 - \lambda_{ij})^{1-y_{ij}}. \quad (38)$$

For the individuals that are not observed, we draw the latent class membership during the missing covariate imputation stage. Recall, because of the complications mentioned in subsection 3.2.2, the Gibbs sampler requires N and \mathcal{X}_{mis} to be sampled simultaneously. In order to implement this stage, the latent class membership, $\boldsymbol{\omega}_i$ of each unobserved individual must be drawn along with, presumably independently, the missing covariates, \mathcal{X}_{mis} . Whereas the distribution from which the covariates is drawn is dependent on the covariate distribution selected (see subsection 4), the ω_i 's are drawn from the discrete distribution with probability vector $\boldsymbol{\phi}_\omega$.

Once again, we use a finite representation of the stick-breaking process for estimation purposes with a sufficiently large upperbound, H_ω^* . Under this construction, the

stick-breaking process is defined as $(\phi_1, \phi_2, \dots, \phi_{H_\omega^*}) = V_h \prod_{h' < h} (1 - V_{h'})$, where $V_1, V_2, \dots, V_{H_\omega^* - 1} \sim \text{Beta}(1, \alpha_\omega)$ and $V_{H_\omega^*} = 1$. To obtain samples from the conditional posterior distribution, $p(\phi_\omega | \dots)$, we first sample V_h by,

$$V_h \sim \text{Beta}\left(1 + N_h, \alpha_\omega + \sum_{h'=h+1}^{H_\omega^*} N_{h'}\right), \quad (39)$$

for all $h = 1, \dots, H_\omega^* - 1$ where N_h is the number of individuals belonging to each latent class. Then compute $\phi_h = V_h \prod_{h' < h} (1 - V_{h'})$ for $h = 1, \dots, H_\omega^*$. The conditional posterior distribution, $p(\alpha_\omega | \dots)$ can be sampled by

$$\alpha_\omega \sim \text{Gamma}\left(a_\omega + H_\omega^* - 1, b_\omega - \sum_{h=1}^{H_\omega^* - 1} \ln(1 - V_h)\right), \quad (40)$$

Lastly, the coefficient vector of latent intercepts, β_ω , can be sampled in the same manner as the other coefficients, β_j .

5.3 Simulations with Unobserved Heterogeneity

The final set of simulations is to demonstrate the importance of accounting for both the observed and unobserved heterogeneity. While the cMLCR model has performed well in the prior sections, it has no method of detecting unobservable heterogeneity and hence produces biased estimation. Interestingly, a model like the LCMCR, which is designed to account for unobservable heterogeneity, performs much better, but still struggles as considerable information can be gained by including the covariates. The BLRCR model when extended to include latent intercepts ($H_\omega > 1$) can utilize the covariates to detect the observable heterogeneity, but also accounts for the additional unobservable heterogeneity. As a result, it is the only model that performs well in this subsection.

Data is simulated for three lists ($J=3$) with one standard normally distributed observed covariate ($H=1$) and one unobserved covariate indicating membership to a latent group with probability 0.35. The coefficients used to simulate the data can be found in [Table 7](#).

Using these coefficients induces positive dependency between lists 1 and 2, but negative dependency between list 3 and the other two lists. If class membership were known, the list probabilities would still be conditionally independent; however, since these covariates are unobserved, the list probabilities are no longer conditionally independent given the observed data.

List (j)	β_{0j}	β_{1j}	$\beta_{\omega j}$
1	-2.5	-1.5	3.0
2	-2.5	-1.5	3.0
3	0.5	-1.5	-3.0

Table 7: Coefficients for Heterogeneity Simulated Data

Table 8 shows the results of 100 simulations on four different population sizes. It becomes obvious that not accounting for unobserved heterogeneity results in biased estimates for both the BLRCR and cMLCR models. Recall, the truth with these simulations is there exists two latent classes. Notice the BLRCR model with the number of hidden classes set at $H_\omega = 2$ performs the best in terms of both accuracy and precision. In a real setting, the number of hidden classes would almost certainly be unknown so we set a sufficiently large value $H_\omega = 20$. While the model doesn't perform quite as well as the aforementioned setting, it dramatically outperforms the methods that do not take unobserved heterogeneity into account.

For situations where observed and unobserved heterogeneity exist, we need a method that accounts for both. It should be noted that the LCMCR model, a methodology that doesn't use covariates, is outperforming the methods that do but assume conditional independence. The LCMCR model is designed to account for unobserved heterogeneity, and since the unobserved heterogeneity plays a substantial role in the capture probability for these simulations, it performs reasonably well. Of course, as we saw in the previous subsections' simulations, the LCMCR does not perform as well as the other methods when most of the heterogeneity can be explained by the covariates.

To further explore the effects of hidden heterogeneity and its detectability, we ran 100 simulations with different coefficients and group percentages. Using the same coefficients in [Table 7](#), but adjusting the absolute value of the coefficients, β_{ω_j} , we created scenarios that depict different strength levels of heterogeneity. Trivially, if the coefficient is 0, there is no unobserved heterogeneity in the capture probabilities. In this situation the algorithm is simply detecting noise that it is mistaking for heterogeneity. On the other hand, when the absolute value of the β_{ω_j} coefficients are set to 5, there is very strong heterogeneity in the capture probabilities for the two groups. The top plot in [Figure 4](#) shows the mean square error (MSE) of the 100 simulations' posterior median of N . A close examination of the plot reveals an overall decrease in the MSE as the heterogeneity strengthens. This is not surprising as the model is attempting to account for heterogeneity, but if it cannot detect the heterogeneity, it will induce bias. When the heterogeneity is stronger, the model is more likely to detect this heterogeneity and account for it properly.

N	Method	N%	MSE	CI Width	CI %
1000	BLRCR($K = 1, H_\omega = 1$)	1.192	209.0	0.386	25.0
	BLRCR($K = 20, H_\omega = 1$)	1.192	207.6	0.382	24.0
	BLRCR($K = 20, H_\omega = 2$)	0.997	57.5	0.334	100.0
	BLRCR($K = 20, H_\omega = 20$)	0.986	67.9	0.385	100.0
	cMLCR	1.231	250.8	0.390	25.0
	Log Linear (BIC)	1.048	94.7	0.199	64.0
	LCMCR	0.851	156.0	0.414	93.0
	Independent	0.875	128.2	0.095	1.0
2000	BLRCR($K = 1, H_\omega = 1$)	1.210	438.8	0.302	1.0
	BLRCR($K = 20, H_\omega = 1$)	1.205	428.4	0.295	1.0
	BLRCR($K = 20, H_\omega = 2$)	1.005	115.4	0.278	99.0
	BLRCR($K = 20, H_\omega = 20$)	0.994	130.0	0.316	98.0
	cMLCR	1.225	469.9	0.274	1.0
	Log Linear (BIC)	1.022	182.7	0.144	44.0
	LCMCR	0.858	301.0	0.388	93.0
	Independent	0.872	259.3	0.067	0.0
5000	BLRCR($K = 1, H_\omega = 1$)	1.229	1161.6	0.223	0.0
	BLRCR($K = 20, H_\omega = 1$)	1.222	1127.8	0.219	0.0
	BLRCR($K = 20, H_\omega = 2$)	1.010	217.2	0.208	100.0
	BLRCR($K = 20, H_\omega = 20$)	1.005	236.1	0.235	100.0
	cMLCR	1.230	1165.8	0.173	0.0
	Log Linear (BIC)	0.954	478.7	0.108	22.0
	LCMCR	0.869	700.0	0.340	85.0
	Independent	0.874	630.3	0.042	0.0
10000	BLRCR($K = 1, H_\omega = 1$)	1.224	2261.2	0.183	0.0
	BLRCR($K = 20, H_\omega = 1$)	1.219	2202.3	0.180	0.0
	BLRCR($K = 20, H_\omega = 2$)	1.005	359.1	0.164	99.0
	BLRCR($K = 20, H_\omega = 20$)	1.008	406.3	0.191	98.0
	cMLCR	1.221	2230.8	0.120	0.0
	Log Linear (BIC)	0.917	935.4	0.088	5.0
	LCMCR	0.869	1426.7	0.322	77.0
	Independent	0.872	1277.3	0.030	0.0

Table 8: Results of 100 capture-recapture simulations per algorithm using a standard normal distribution for the known covariate and 0.35 probability of belonging to the latent class.




Figure 4: Effect of Strength of Heterogeneity on the Mean Square Error

The bottom plot in Figure 4 shows the MSE of the average posterior group identifier for the observed data. In other words, if the first observation actually belongs to the hidden group, $\omega_1=1$, then we find the average number of times the algorithm placed the observation in the first group, $\bar{\omega}_1$. In the binary case, $H_\omega = 2$, the MSE can be computed as

$$MSE = \frac{1}{n} \sum_{i=1}^n (\bar{\omega}_i - \omega_i)^2. \quad (41)$$

As the heterogeneity strengthens, the detectability of the unobserved groupings increases and the MSE drops considerably. In the case where the absolute value of $\beta_{\omega j}$ is set to 5, in many of the simulations between 90 and 95% of the observations are detected correctly in

over 90% of posterior samples.

Unobserved heterogeneity creates a trade off. On one hand, if the unobserved heterogeneity is weak, there will be relatively little bias in N ; however, the ability to detect this heterogeneity decreases as well. On the other hand, When the unobserved heterogeneity is strong, the bias in N will be relatively larger, but the ability to detect and properly account for this heterogeneity increases.

6 Model Selection

The methodology described in the previous sections assumes the practitioner knows exactly how the covariates impact the probability of each individual going undetected on a given list. The key to the successful estimation of the population size relies on the proper and accurate sampling of the coefficients in the logistic regression, or at a minimum, these coefficients must lead to a correct estimate of the overall probability of an individual being missing. Hence, one must carefully select which covariates to use, how to use these covariates, and whether there are any interactions that may play an important role in this process. While it may be beneficial to rely on domain knowledge of the population to make these determinations, this may not be sufficient so it is desirable to take a statistical approach to this problem.

6.1 Model Comparison via the Bayes Factor

Perhaps the most ubiquitous approach to model selection in a Bayesian setting is through the computation of the Bayes Factor. This approach of model comparison has a lot of similarities with the problem of hypothesis testing. Suppose there is a null and alternative model, \mathcal{M}_0 and \mathcal{M}_1 , respectively. With the standard frequentist approach, the practitioner computes the the probability of observing the data given the null model, the p-value, and chooses to accept or reject the null model based on some arbitrary threshold, typically

based on controlling for Type I error. On the other hand, using a Bayesian approach (Jeffreys, 1935, 1967), we compare two hypotheses or models by directly computing the probability of each model given the data through the posterior probabilities. This means, unlike frequentist hypothesis testing, there is no need to distinguish between which model is the null and which is the alternative.

For the BLRCR model, the posterior probability of \mathcal{M}_1 given the observed data is written as

$$p(\mathcal{M}_1|\mathcal{Y}_{obs}, \mathcal{X}_{obs}) = \frac{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|\mathcal{M}_1)p(\mathcal{M}_1)}{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs})}. \quad (42)$$

Next, we calculate the posterior odds ratio between \mathcal{M}_1 and \mathcal{M}_0 as

$$\frac{p(\mathcal{M}_1|\mathcal{Y}_{obs}, \mathcal{X}_{obs})}{p(\mathcal{M}_0|\mathcal{Y}_{obs}, \mathcal{X}_{obs})} = \frac{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|\mathcal{M}_0)p(\mathcal{M}_0)}{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|\mathcal{M}_1)p(\mathcal{M}_1)}. \quad (43)$$

If we do not a priori favor one model over another, such that $p(\mathcal{M}_0) = p(\mathcal{M}_1)$, then the posterior odds ratio is equal to what is known as the Bayes Factor,

$$BF_{10} = \frac{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|\mathcal{M}_1)}{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|\mathcal{M}_0)}. \quad (44)$$

where values greater than 1 would suggest \mathcal{M}_1 , and values less than 1 would suggest \mathcal{M}_0 as the better fit.

Calculation of the Bayes Factor is often difficult or intractable as it is the ratio of the marginal likelihoods of the two models. Given a specific model, \mathcal{M}_i , The marginal likelihood is the probability of the data integrated across all values of the parameters,

$$p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|\mathcal{M}_i) = \int p(\mathcal{Y}_{obs}, \mathcal{X}_{obs}|\boldsymbol{\theta}, \mathcal{M}_i)p(\boldsymbol{\theta}|\mathcal{M}_i)d\boldsymbol{\theta}. \quad (45)$$

This is the the denominator in the calculation of the posterior distribution assuming a fixed model (ignore the \mathcal{M}_i). Because this is just the normalizing constant on the posterior density and because of the complexity of its calculation, the calculation of the marginal likelihood is often avoided by taking advantage of posterior/prior conjugacy or normalization in the discrete case.

Without the ability to solve Equation 45 analytically, we must rely on an approximate solution. Perhaps the simplest solution is to recognize Equation 45 can be written as the expected value of the likelihood under the prior distribution. As such, a simple application of the Law of Large numbers, leads to the estimator,

$$\hat{p}_1(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \mathcal{M}_i) = \frac{1}{G} \sum_{g=1}^G p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta}^{(g)}, \mathcal{M}_i), \quad (46)$$

where G is the number of samples and $\boldsymbol{\theta}^{(g)}$ are the sampled parameter values from the **prior** distribution. While this estimator is simple and easy to implement, it converges slowly when the prior and posterior are significantly different, which is often the case when using non-informative priors. Many of the prior samples will produce negligibly small likelihood values, leading to a large variance in the estimator. One solution to improve this integration is to use importance sampling Kass and Raftery (1995); Robert et al. (2009), but leads to the additional hardship of selecting a proper importance sampling function. Furthermore, this method requires sampling from the prior which creates additional difficulties when using an improper prior such as the prior for N as with the BLRCR. Instead of using the prior distribution, Newton and Raftery (1994) showed that the harmonic mean can be used to approximate the marginal likelihood using posterior samples,

$$\hat{p}_2(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \mathcal{M}_i) = \left[\frac{1}{G} \sum_{g=1}^G \frac{1}{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta}^{(g)}, \mathcal{M}_i)} \right]^{-1}, \quad (47)$$

where $\boldsymbol{\theta}^{(g)}$ are the sampled parameter values from the **posterior** distribution.

Unfortunately, this estimator tends to not work well with results being unstable and the estimator having infinite variance (Neal, 1994; Robert et al., 2009). To improve this estimator, Newton and Raftery (1994) present additional estimators using mixtures from the prior and posterior. Alternatively, Gelfand and Dey (1994) derives an estimator in the context of importance sampling where the posterior distribution is selected as the importance sampling density.

$$\begin{aligned} E_{p(\boldsymbol{\theta} | \mathcal{Y}_{obs}, \mathcal{X}_{obs})} \left[\frac{\psi(\boldsymbol{\theta})}{p(\boldsymbol{\theta})p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta})} \right] &= \int \frac{\psi(\boldsymbol{\theta})}{p(\boldsymbol{\theta})p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta})} \cdot p(\boldsymbol{\theta} | \mathcal{Y}_{obs}, \mathcal{X}_{obs}) d\boldsymbol{\theta} \\ &= \int \frac{\psi(\boldsymbol{\theta})}{p(\boldsymbol{\theta})p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta})} \cdot \frac{p(\boldsymbol{\theta})p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta})}{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs})} d\boldsymbol{\theta} \\ &= \frac{1}{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs})} \int \psi(\boldsymbol{\theta}) d\boldsymbol{\theta} \\ &= \frac{1}{p(\mathcal{Y}_{obs}, \mathcal{X}_{obs})} \end{aligned} \quad (48)$$

where $\psi(\boldsymbol{\theta})$ can be any proper density. Hence a consistent estimator for the marginal likelihood is

$$\hat{p}_3(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \mathcal{M}_i) = \left[\frac{1}{G} \sum_{g=1}^G \frac{\psi(\boldsymbol{\theta}^{(g)})}{p(\boldsymbol{\theta}^{(g)} | \mathcal{M}_i)p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta}^{(g)}, \mathcal{M}_i)} \right]^{-1}. \quad (49)$$

Notice that if we set $\psi(\boldsymbol{\theta}^{(g)}) = p(\boldsymbol{\theta}^{(g)})$, we have the original harmonic mean estimator, $\hat{p}_2(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \mathcal{M}_i)$. While any density can be selected, we need a density with lighter tails than $p(\boldsymbol{\theta}^{(g)} | \mathcal{M}_i)p(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \boldsymbol{\theta}^{(g)}, \mathcal{M}_i)$ in order to improve the original estimator and achieve

finite variance (Robert et al., 2009). Potential candidate densities that may work include asymptotic densities around the maximum likelihood values or kernel approximations. For example, in the case of Probit regression, Marin and Robert (2018) found success using the normal density with mean and covariance equal to the maximum likelihood estimate.

This subsection provides a handful of methods for computing the Bayes Factor for the BLRCR model, but there are several potential issues making the above procedures impossible or at best questionable to implement. For starters, the use of improper priors on N means the Bayes Factors should certainly not be interpreted literally as the actual value is off by an arbitrary constant. Second, the multi-dimensionality of the problem leads to issues of converge with even the largest of chains. Nevertheless, we will see in [INSERT SIMULATION SUBSECTION] that the Bayes Factor computations above lead us to the correct implementation more often than not.

Our focus is on the methods that utilize samples from the posterior distribution since the algorithm provided in subsection 3.2.2 utilizes a MCMC method. Further work could be done on methods that use samples from the prior like the estimator, $\hat{p}_1(\mathcal{Y}_{obs}, \mathcal{X}_{obs} | \mathcal{M}_i)$. In addition, other approaches such as asymptotic approximations like Laplace’s Method or exploiting functional equalities Chib (1995) may prove promising as well. While we leave this to future research, we believe a more promising approach may be through the use of regression regularization.

6.2 Regularization and Prior Specifications

We propose a different approach to model selection via the Bayes factor by instead narrowing in on the portion of the methodology that samples the logistic regression coefficients. While the various subsections have added stages to the BLRCR sampler to accommodate various complications, the sampling stages can be summarized into two basic groupings as presented in subsection 3.2.2. The first grouping of sampling stages are used

to augment the missing data by first sampling the size of the data, N , and then sampling the missing covariates, \mathcal{X}_{mis} . The second grouping of sampling stages then samples the probability an individual goes missing through a deterministic function of the sampled logistic regression coefficients, β . Under this framework, taking a sample of regression coefficients using the augmented data is no different than taking a sample of regression coefficients from a dataset that is static. Hence, we can narrow our focus to finding the correct fitting for the linear component of our logistic regression.

By narrowing our focus to the regression coefficients, we can include all potential variables (including transformations and interactions) with the purported threat of overfitting.

There are multiple ways to overcome the problem with overfitting. A typical way seen especially in the data science community is through cross validation using training and testing sets. This is not a feasible approach for the capture-recapture problem, as the number of observed captures, plays a crucial rule in the estimation process. Theoretically, there may be ways around this by estimating random subpopulation sizes, but as is typical with these approaches, less data being used for estimation leads to more uncertainty especially in the face of sparsity. As a result, we suggest implementing regularized regression, which in a Bayesian context, is through different types of priors.

[MENTION HORSESHOE and SLAB/SPIKE PRIOR HERE ALONG WITH PAPERS FOR IMPLEMENTING, note that using normal priors is a form of regularization, i.e. ridge regression]

While regularization has its shortcomings, using these different priors avoids many of the complications that can arise with the computation of the Bayes Factor. Also, when the number of model comparisons is large, computing the Bayes Factor requires posterior samples from each of the candidate models, whereas with regularization, only one MCMC chain is required.

7 Example 1 (probably superhero)

8 Example 2 (find something real even if we don't know the answer)

9 Conclusion

The objective of this dissertation project is to develop a Bayesian capture-recapture (CR) method that can utilize covariate information to understand the heterogeneity between individuals. Presented in this proposal is a framework for modelling capture-recapture with covariates. Specifically, we develop the Bayesian Logistic Regression Capture Recapture (BLRCR) that utilizes the covariate information directly in estimating individual capture rates. Further, we account for unobserved heterogeneity with the use of latent classes. This model is solved using an MCMC algorithm to approximate the posterior distribution of our population size, N . I propose three primary areas of interest that will be addressed as the focus of the dissertation: unobserved heterogeneity, modelling the covariate distribution, and covariate selection.

9.1 Unobserved Heterogeneity

This extension concerns the issue of unobserved heterogeneity described in subsection 5. To account for this heterogeneity, the algorithm detects latent groups and applies different intercepts in the linear term that defines the capture probability. Currently the number of latent classes is fixed and must be selected by the practitioner. An extension would be to include some method for determining the number of latent classes to include. While it may be reasonable in some cases to know the number of latent classes, an obvious approach would be to incorporate a stick-breaking prior on the number of latent classes.

9.2 Modelling the Covariate Distribution

A requirement of the algorithm is the necessity of specifying a distribution for the missing covariates, discussed in subsection 4. Typically this distribution is unknown so utilizing a distribution that is fairly flexible is ideal. In this proposal, we implemented the non-parametric approach of using an infinite mixture of normals with membership

determined through a stick-breaking process (and additionally a single multivariate normal which is a special case). While the infinite mixture of normal distributions tends to perform well when the distribution is at least somewhat normal, it struggles when presented with certain non-normal distributions. Alternatively, I propose using Dirichlet process mixtures as described in [Gelman et al. \(2014\)](#). A sufficiently large concentration parameter, $\alpha \rightarrow \infty$, would be in essence the same thing as the empirical distribution.

9.3 Covariate Selection

If many covariates are present, we need a method for determining which covariates should be used while also considering interaction effects. Along these same lines, we need to consider how a variable affects the capture probability as some sort of transformation may be required. We suggest three possible routes including model averaging, variable selection through priors, or even post-hoc analysis like computing the Bayes factor.

10 Appendix A: Conditional Maximum Likelihood Estimation

Instead of using the full likelihood we replace it with the conditional likelihood, conditioned on each individual being observed once.

$$\begin{aligned} p(\mathcal{Y}|N, \boldsymbol{\beta}, \mathcal{X}_{obs}, \mathcal{X}_{mis}; \mathbf{y} \neq \mathbf{0}) &= \frac{p(\mathcal{Y}|N, \boldsymbol{\beta}, \mathcal{X}_{obs}, \mathcal{X}_{mis})}{P(\mathbf{y} \neq \mathbf{0})} \\ &= \prod_{i=1}^N \frac{\prod_{j=1}^J \lambda_{ij}^{y_{ij}} (1 - \lambda_{ij})^{1-y_{ij}}}{1 - \prod_{j=1}^J (1 - \lambda_{ij})} \end{aligned} \quad (50)$$

The posterior distribution can therefore be written as

$$\begin{aligned} p(N, \boldsymbol{\beta}, \mathcal{X}_{mis} | \mathcal{Y}, \mathcal{X}_{obs}) &\propto \left[\prod_{i=1}^N \frac{\prod_{j=1}^J \lambda_{ij}^{y_{ij}} (1 - \lambda_{ij})^{1-y_{ij}}}{1 - \prod_{j=1}^J (1 - \lambda_{ij})} \right] \times \left[\prod_{i=1}^N \mathbf{g}(\boldsymbol{\theta}_i) \right] \times \left[\frac{1}{N} \right] \\ &\times \left[\prod_{j=1}^J \left(\frac{1}{2\pi} \right)^{H/2} |B|^{-1/2} e^{-\frac{1}{2}(\mathbf{b} - \boldsymbol{\beta}_j)^T B^{-1}(\mathbf{b} - \boldsymbol{\beta}_j)} \right]. \end{aligned} \quad (51)$$

Instead of computing a MCMC sampler, we will instead compute the Maximum a posteriori (MAP) estimate for the coefficients, $\boldsymbol{\beta}_{MAP}$, using gradient ascent. Then, this estimate can be plugged into a Horvitz-Thompson estimator to find the N_{MAP} .

Taking into consideration only the parts of the posterior that depend on $\boldsymbol{\beta}$, the log posterior is

$$\begin{aligned} \ln p(N, \boldsymbol{\beta}, \mathcal{X}_{mis} | \mathcal{Y}, \mathcal{X}_{obs}) &\propto -\frac{1}{2} \sum_{j=1}^J (\mathbf{b} - \boldsymbol{\beta}_j)^T B^{-1}(\mathbf{b} - \boldsymbol{\beta}_j) \\ &+ \sum_{i=1}^N \sum_{j=1}^J y_{ij} \ln(\lambda_{ij}) + (1 - y_{ij}) \ln(1 - \lambda_{ij}) \\ &+ \ln(1 - \prod_{j=1}^J (1 - \lambda_{ij})) \end{aligned} \quad (52)$$

Taking the first derivative with respect to β_j gives the gradient

$$\frac{\Delta \ln p(N, \beta, \mathcal{X}_{mis} | \mathcal{Y}, \mathcal{X}_{obs})}{\Delta \beta_j} = B^{-1}(b - \beta_j) + \sum_{i=1}^n \left(y_{ij} - \lambda_{ij} + \frac{\lambda_{ij} \prod_{j=1}^J (1 - \lambda_{ij})}{1 - \prod_{j=1}^J (1 - \lambda_{ij})} \right) \mathbf{x}_i^T. \quad (53)$$

Notice in the equation above that only the observed data appears. This convenience occurs due to the fact that $E[y_{ij} | \mathbf{y}_i = \mathbf{0}] = 0$ for all individuals ([Alho, 1990](#)). Using the gradient, apply the gradient ascent algorithm until convergence to obtain β_{MAP} .

11 Appendix B: LCMCR Model

This section summarizes the Bayesian Non-Parametric Latent-Class Capture-Recapture (LCMCR) derivation and explanations provided in [Manrique-Vallier \(2016\)](#). The LCMCR model is a framework for the capture-recapture (CR) problem that uses the Bayesian nonparametric latent-class model (NPLCM) proposed in [Dunson and Xing \(2009\)](#) to model $f(\mathbf{y}_i|\theta)$.

In order to account for unobserved heterogeneity, each individual, i , is modeled such that they belong to a hidden, latent class, z_i , with probability π_k . After determining a latent class, an individual is captured according to a Bernoulli distribution on list j with probability, λ_{jk} . This is known as the latent-class model ([Goodman, 1974](#)), and yields the probability mass function

$$f(\mathbf{y}_i|\boldsymbol{\lambda}, \boldsymbol{\pi}) = \sum_{k=1}^K \pi_k \prod_{j=1}^J \lambda_{jk}^{y_{ij}} (1 - \lambda_{jk})^{1-y_{ij}}, \quad (54)$$

where $\boldsymbol{\pi} = (\pi_1, \dots, \pi_K)$ and $\boldsymbol{\lambda} = (\lambda_{jk})$ for $j = 1, \dots, J$ and $k = 1, \dots, K$.

Inserting [Equation 54](#) into [Equation 1](#) yields the following likelihood equation

$$p(\mathcal{Y}|\boldsymbol{\lambda}, \boldsymbol{\pi}, N) = \binom{N}{n} \left[\sum_{k=1}^K \pi_k \prod_{j=1}^J (1 - \lambda_{jk}) \right]^{N-n} \times \prod_{i=1}^n \sum_{k=1}^K \pi_k \prod_{j=1}^J \lambda_{jk}^{y_{ij}} (1 - \lambda_{jk})^{1-y_{ij}}. \quad (55)$$

The number of latent classes is endogenized with the probability of belonging to each latent class, π_k , being drawn from a “stick-breaking” process ([Sethuraman, 1991](#)). The parameter, α , controls the amount of concentration of the probability mass. In other words, larger values of α will lead to a larger number of relevant latent classes. A Beta(1,1) prior distribution is placed on each of the $J \times K^*$ probabilities, and a Gamma(a, b) prior is placed on α .

The model can be summarized through the following hierarchical generative process

$$\begin{aligned}
y_{ij}|z_i &\sim \text{Bernoulli}(\lambda_{jz}) \quad \text{for } j = 1, \dots, J \text{ and } i = 1, \dots, N \\
z_i &\sim \text{Discrete}(\{1, 2, \dots\}, (\pi_1, \pi_2, \dots)) \quad \text{for } i = 1, \dots, N \\
\lambda_{jk} &\sim \text{Beta}(1, 1) \quad \text{for } j = 1, \dots, J \text{ and } k = 1, 2, \dots \\
(\pi_1, \pi_2, \dots) &\sim \text{SB}(\alpha) \\
\alpha &\sim \text{Gamma}(a, b).
\end{aligned} \tag{56}$$

12 References

- Agresti, A. (1994). Simple Capture-Recapture Models Permitting Unequal Catchability and Variable Sampling Effort. *Biometrics*, 50(2):494–500.
- Alho, J. M. (1990). Logistic Regression in Capture-Recapture Models. *Biometrics*, 46(3):623.
- Alpizar-Jara, R. and Pollock, K. H. (1996). A combination line transect and capture-recapture sampling model for multiple observers in aerial surveys. *Environmental and Ecological Statistics*, 3(4):311–327.
- Baker, S. G. (1990). A Simple EM Algorithm for Capture-Recapture Data with Categorical Covariates. *Biometrics*, 46(4):1193.
- Ball, P., Asher, J., Sulmont, D., and Manrique, D. (2003). An estimate of the total number of victims killed or disappeared in the armed internal conflict between 1980 and 2000. *AAAS. Report to the Peruvian Truth and Reconciliation Commission (CVR)*.
- Ball, P., Betts, W., Scheuren, F., Dudukovich, J., and Asher, J. (2002). Killings and Refugee Flow in Kosovo, March - June 1999: A Report to the International Criminal Tribunal for the Former Yugoslavia. *American Association for the Advancement of Science (AAAS)*.
- Ball, P. and Price, M. (2018). The Statistics of Genocide. *CHANCE*, 31(1):38–45.
- Ball, P., Spierer, H. F., and Spierer, L. (2000). *Making the case: investigating large scale human rights violations using information systems and data analysis*. American Association for the Advancement of Science, Washington, D.C.
- Basu, S. (2001). Bayesian capture-recapture methods for error detection and estimation of population size: Heterogeneity and dependence. *Biometrika*, 88(1):269–279.

- Bonner, S. and Schofield, M. (2014). MC(MC)MC: exploring Monte Carlo integration within MCMC for mark-recapture models with individual covariates. *Methods in Ecology and Evolution*, 5(12):1305–1315.
- Borchers, D. L., Zucchini, W., and Fewster, R. M. (1998). Mark-Recapture Models for Line Transect Surveys. *Biometrics*, 54(4):1207–1220.
- Castledine, B. J. (1981). A Bayesian Analysis of Multiple-Recapture Sampling for a Closed Population. 68:197–210.
- Chapman, D. G. (1954). The Estimation of Biological Populations. *The Annals of Mathematical Statistics*, 25(1):1–15.
- Chib, S. (1995). Marginal Likelihood from the Gibbs Output. *Journal of the American Statistical Association*, 90(432):1313–1321.
- Darroch, J. N. (1958). The Multiple-Recapture Census: I. Estimation of a Closed Population. *Biometrika*, 45(3/4):343–359.
- Darroch, J. N., Fienberg, S. E., Glonek, G. F. V., and Junker, B. W. (1993). A Three-Sample Multiple-Recapture Approach to Census Population Estimation With Heterogeneous catchability. *Journal of the American Statistical Association*, page 13.
- Dunson, D. B. and Park, J.-H. (2008). Kernel stick-breaking processes. *Biometrika*, 95(2):307–323.
- Dunson, D. B. and Xing, C. (2009). Nonparametric Bayes Modeling of Multivariate Categorical Data. *Journal of the American Statistical Association*, 104(487):1042–1051.
- Fienberg, S. E. (1972). The Multiple Recapture Census for Closed Populations and Incomplete 2k Contingency Tables. *Biometrika*, 59(3):591–603.

- Fienberg, S. E., Johnson, M. S., and Junker, B. W. (1999). Classical Multilevel and Bayesian Approaches to Population Size Estimation Using Multiple Lists. *Journal of Royal Statistical Society*, 162(3):383–405.
- Geiger, H. and Werner, A. (1924). Die Zahl der von Radium ausgesandten α -Teilchen. *Zeitschrift für Physik*, 21:187–201.
- Gelfand, A. E. and Dey, D. K. (1994). Bayesian Model Choice: Asymptotics and Exact Calculations. *Journal of the Royal Statistical Society: Series B (Methodological)*, 56(3):501–214.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., and Rubin, D. B. (2014). *Bayesian Data Analysis*. Chapman & Hall/CRC, third edition.
- George, E. I. (1992). Capture-Recapture Estimation Via Gibbs Sampling. *Biometrika*, 79(4):677–683.
- Goldberg, J. D. and Wittes, J. T. (1978). The Estimation of False Negatives in Medical Screening. *Biometrics*, 34(1):77–86.
- Goodman, L. A. (1974). Exploratory Latent Structure Analysis Using Both Identifiable and Unidentifiable Models. *Biometrika*, 61(2):215–231.
- Guberek, T., Guzman, D., Price, M., Lum, K., and Ball, P. (2010). To Count the Uncounted: An Estimation of Lethal Violence in Casanare. *Benetech Human Rights Program*, pages 1–31.
- Henderson, P. A. and Southwood, T. (2016). *Ecological Methods*. John Wiley and Sons, 4 edition.
- Horvitz, D. G. and Thompson, D. J. (1952). A Generalization of Sampling Without Replacement From a Finite Universe. *Journal of the American Statistical Association*, 47(260):663–685.

- Huggins, R. M. (1989). On the Statistical Analysis of Capture Experiments. *Biometrika*, 76(1):133–140.
- Ishwaran, H. and James, L. F. (2001). Gibbs Sampling Methods for Stick-Breaking Priors. *Journal of the American Statistical Association*, 96(453):161–173.
- Jeffreys, H. (1935). Some Tests of Significance, Treated by the Theory of Probability. *Mathematical Proceedings of the Cambridge Philosophical Society*, 31(2):203–222.
- Jeffreys, H. (1967). *Theory of Probability*. Clarendon Press, Oxford, 3rd edition.
- Kass, R. E. and Raftery, A. E. (1995). Bayes Factors. *Journal of the American Statistical Association*, 90(430):773–795.
- Khabsa, M. and Giles, C. L. (2014). The Number of Scholarly Documents on the Public Web. *PLoS ONE*, 9(5):e93949.
- King, R. and Brooks, S. P. (2008). On the Bayesian Estimation of a Closed Population Size in the Presence of Heterogeneity and Model Uncertainty. *Biometrics*, 64(3):816–824.
- King, R., McClintock, B. T., Kidney, D., and Borchers, D. (2016). Capture–recapture abundance estimation using a semi-complete data likelihood approach. *The Annals of Applied Statistics*, 10(1):264–285.
- Lawrence, S. and Giles, C. L. (1998). Searching the World Wide Web. *Science*, 280(5360):98–100.
- Lincoln, F. C. (1930). Calculating Waterfowl Abundance on the Basis of Banding Returns. *United States Department of Agriculture*, Circular 118.
- Manrique-Vallier, D., Ball, P., and Sadinle, M. (2020). Capture-Recapture for Casualty Estimation and Beyond: Recent Advances and Research Directions.

- Manrique-Vallier, D., Ball, P., and Sulmont, D. (2019). Estimating the Number of Fatal Victims of the Peruvian Internal Armed Conflict, 1980-2000: an application of modern multi-list Capture-Recapture techniques.
- Manrique-Vallier, D. and Fienberg, S. E. (2008). Population Size Estimation Using Individual Level Mixture Models. *Biometrical Journal*, 50(6):1051–1063.
- Manrique-Vallier, D. (2016). Bayesian Population Size Estimation using Dirichlet Process Mixtures. *Biometrics*, 72(4):1246–1254.
- Marin, J.-M. and Robert, C. P. (2018). Importance sampling methods for Bayesian discrimination between embedded models. arXiv:0910.2325 [stat].
- Neal, R. (1994). Contribution to the discussion of Approximate Bayesian inference with the weighted likelihood bootstrap" by Michael A. Newton and Adrian E. Raftery. *Journal of Royal Statistical Society Series B*, 56(1):41–42.
- Newton, M. A. and Raftery, A. E. (1994). Approximate Bayesian Inference with the Weighted Likelihood Bootstrap. *Journal of the Royal Statistical Society: Series B (Methodological)*, 56(1):3–26.
- Otis, D. L., Burnham, K. P., White, G. C., and Anderson, D. R. (1978). Statistical Inference from Capture Data on Closed Animal Populations. *Wildlife Monographs*, (62,):3–135.
- Petersen, C. J. (1895). The yearly immigration of young plaice into the limfjord from the german sea. *The Danish Biological Station*.
- Pledger, S. (2000). Unified Maximum Likelihood Estimates for Closed Capture-Recapture Models Using Mixtures. *Biometrics*, 56(2):434–442.
- Pollock, K. H. (2002). The use of auxiliary variables in capture-recapture modelling: An overview. *Journal of Applied Statistics*, 29(1-4):85–102.

- Pollock, K. H., Hines, J. E., and Nichols, J. D. (1984). The Use of Auxiliary Variables in Capture-Recapture and Removal Experiments. *Biometrics*, 40(2):329–340.
- Polson, N. G., Scott, J. G., and Windle, J. (2013). Bayesian inference for logistic models using Polya-Gamma latent variables. *Journal of the American Statistical Association*, 108(504):1339–1349.
- Robert, C. P., Wraith, D., Goggans, P. M., and Chan, C.-Y. (2009). Computational methods for Bayesian model choice. pages 251–262, Oxford (Mississippi).
- Roberts, H. V. (1967). Informative Stopping Rules and Inferences about Population Size. *Journal of the American Statistical Association*, 62:763–775.
- Royle, J. A. (2009). Analysis of Capture-Recapture Models with Individual Covariates Using Data Augmentation. *Biometrics*, 65(1):267–274.
- Royle, J. A., Dorazio, R. M., and Link, W. A. (2007). Analysis of Multinomial Models With Unknown Index Using Data Augmentation. *Journal of Computational and Graphical Statistics*, 16(1):67–85.
- Rubin, D. B. (1976). Inference and Missing Data. *Biometrika*, 63(3):581–592.
- Rubin, D. B. (1981). The Bayesian Bootstrap. *The Annals of Statistics*, 9(1):130–134.
- Sanathanan, L. (1972). Models and Estimation Methods in Visual Scanning Experiments. *Technometrics*, 14(4):813–829.
- Sandland, R. L. and Cormack, R. M. (1984). Statistical Inference for Poisson and Multinomial Models for Capture- Recapture Experiments. *Biometrika*, 71(1):27–33.
- Schaefer, M. B. (1951). Estimation of Size of Animal Populations by Marking Experiments. *Fishery Bulletin*, 52(69):191–203.

- Schnabel, Z. E. (1938). The Estimation of Total Fish Population of a Lake. *The American Mathematical Monthly*, 45(6):348.
- Sekar, C. C. and Deming, E. (1949). On a Method of Estimating Birth and Death Rates and the Extent of Registration. *Journal of the American Statistical Association*, 44(245):101–115.
- Sethuraman, J. (1991). A Constructive Definition of Dirichlet Priors:. Technical report, Defense Technical Information Center, Fort Belvoir, VA.
- Smith, P. J. (1991). Bayesian Analyses for a Multiple Capture-Recapture Model. *Biometrika*, 78(2):399–407.
- Stoklosa, J., Hwang, W.-H., Wu, S.-H., and Huggins, R. (2011). Heterogeneous Capture-Recapture Models with Covariates: A Partial Likelihood Approach for Closed Populations. *Biometrics*, 67(4):1659–1665.
- Tilling, K. and Sterne, J. A. C. (1999). Capture-Recapture Models Including Covariate Effects. *American Journal of Epidemiology*, 149(4):392–400.
- Woodbury, M., Clive, J., and Garson, A. (1978). Mathematical typology: a grade of membership technique for obtaining disease definition. *Computers and Biomedical Research*, 11(3):277–298.
- Yee, T. W., Stoklosa, J., and Huggins, R. M. (2015). The **VGAM** Package for Capture-Recapture Data Using the Conditional Likelihood. *Journal of Statistical Software*, 65(5).
- Zwane, E. and van der Heijden, P. (2003). Implementing the parametric bootstrap in capture–recapture models with continuous covariates. *Statistics & Probability Letters*, 65(2):121–125.

Zwane, E. and van der Heijden, P. (2005). Population estimation using the multiple system estimator in the presence of continuous covariates. *Statistical Modelling*, 5(1):39–52.

Curriculum Vitae